

High-Level UDP / TCP/IP

Network Programming

Overview

- UDP (User Datagram Protocol) is unreliable
 - Potentially delivered out of order (or not at all!)
 - Connectionless
- TCP (Transmission Control Protocol) provides reliability
 - In-order
 - Connection-oriented
- Both sit on top of another protocol

Internet Protocol (IP)

- IP is the network layer
- Essentially, responsible for host to host packet delivery (routing)
- Translation between multiple data link protocols such as ethernet, wifi, etc.

IP Datagrams

- IP provides connectionless, unreliable delivery of IP datagrams
- Connectionless: All datagrams are independent of each other
- Unreliable: no guarantee datagrams are even delivered, let alone ordered

IP Addresses

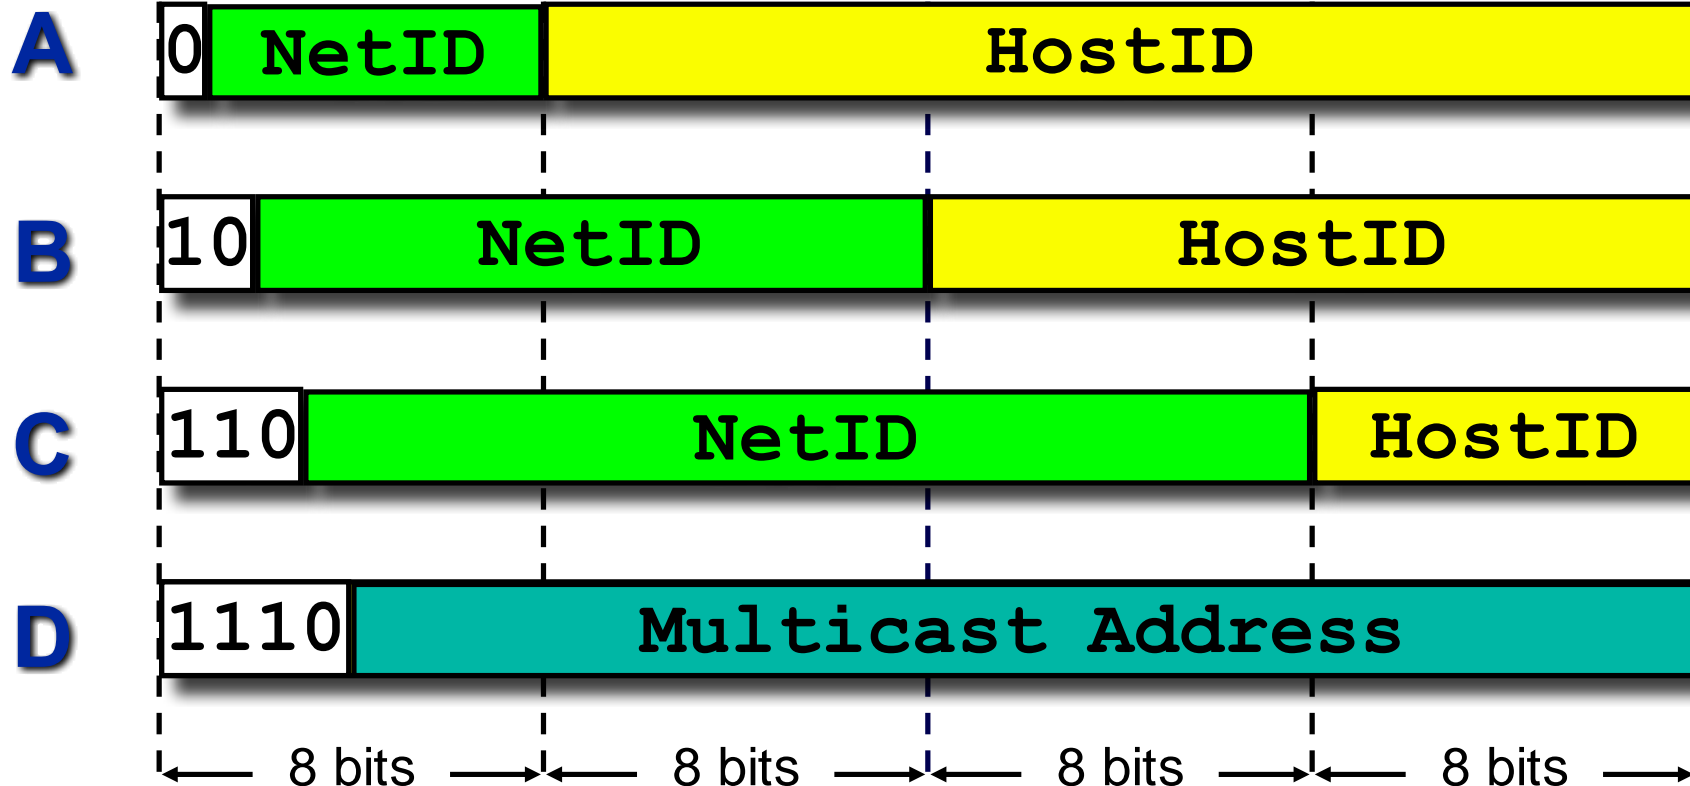
- IP at network layer but must be able to talk to other devices on different mediums! (ex: iPhone to wired server)
 - Also why MAC at different layer than IP
- Must provide some degree of *network* information
 - Allows for efficient routing of datagrams

IP Addresses

- IP addresses are logical addresses
- Four octets (32 bits) [IPv4]
- Includes a network portion and a host portion
- All IPs must be unique

IP Address Format

Class



Ramifications

- Class A: 128 network IDs
 - 16M host IDs per network ID
- Class B: 16K network IDs
 - 64K host IDs per network ID
- Class C: 2M network IDs
 - 256 host IDs per network ID

Network / Host IDs

- Network IDs are assigned by a central authority
 - ICANN, IANA
- Host IDs are assigned locally by a systems administrator
- Network ID and host ID are used for routing purposes

IP Address Format

- IPs are often written in *dotted decimal* notation
- For example, 128.113.0.2 (www.rpi.edu)
 - 10000000.01110001.00000000.00000010
- RPI must have a class B address!
 - Leading digits are 10

Network / Host Addresses

- Hosts aren't assigned addresses, their network interface is
- Hosts may have multiple NICs
- If the network addresses are the same, they share the network

Broadcast / Network Addresses


- Broadcast address: host ID all 1's
- Broadcasts may be implemented however the underlying layer sees fit
- Network address: host ID all 0's, refers to entire network

Subnet Addresses

- Organizations are able to further subdivide its available address space into “subnets”
- For example, clump nearby machines into their own subnet
 - Could also do this logically




Subnetting



Subnetting

- IP subnet broadcasts have host ID of all 1s
- Subnets can simplify routing
- Multiple subnets can share a single wire!

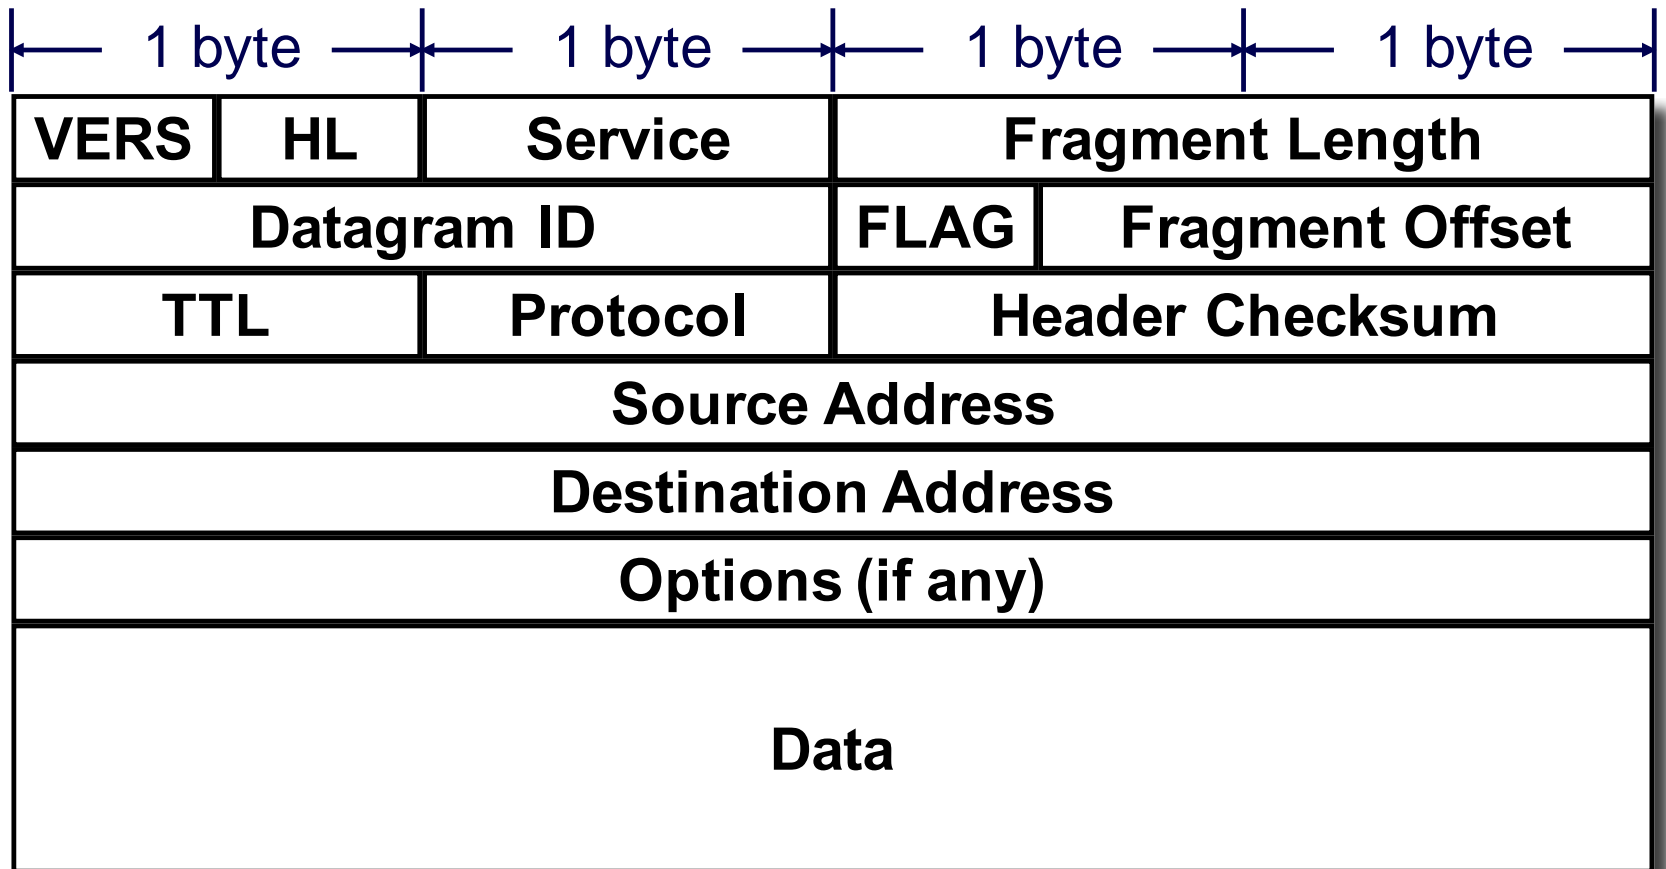


IP

- Connectionless delivery
 - Each datagram handled individually
- Unreliable
 - No guarantee
- Fragmentation + reassembly
 - Hardware MTU
- Routing
- Error detection

IP Datagram

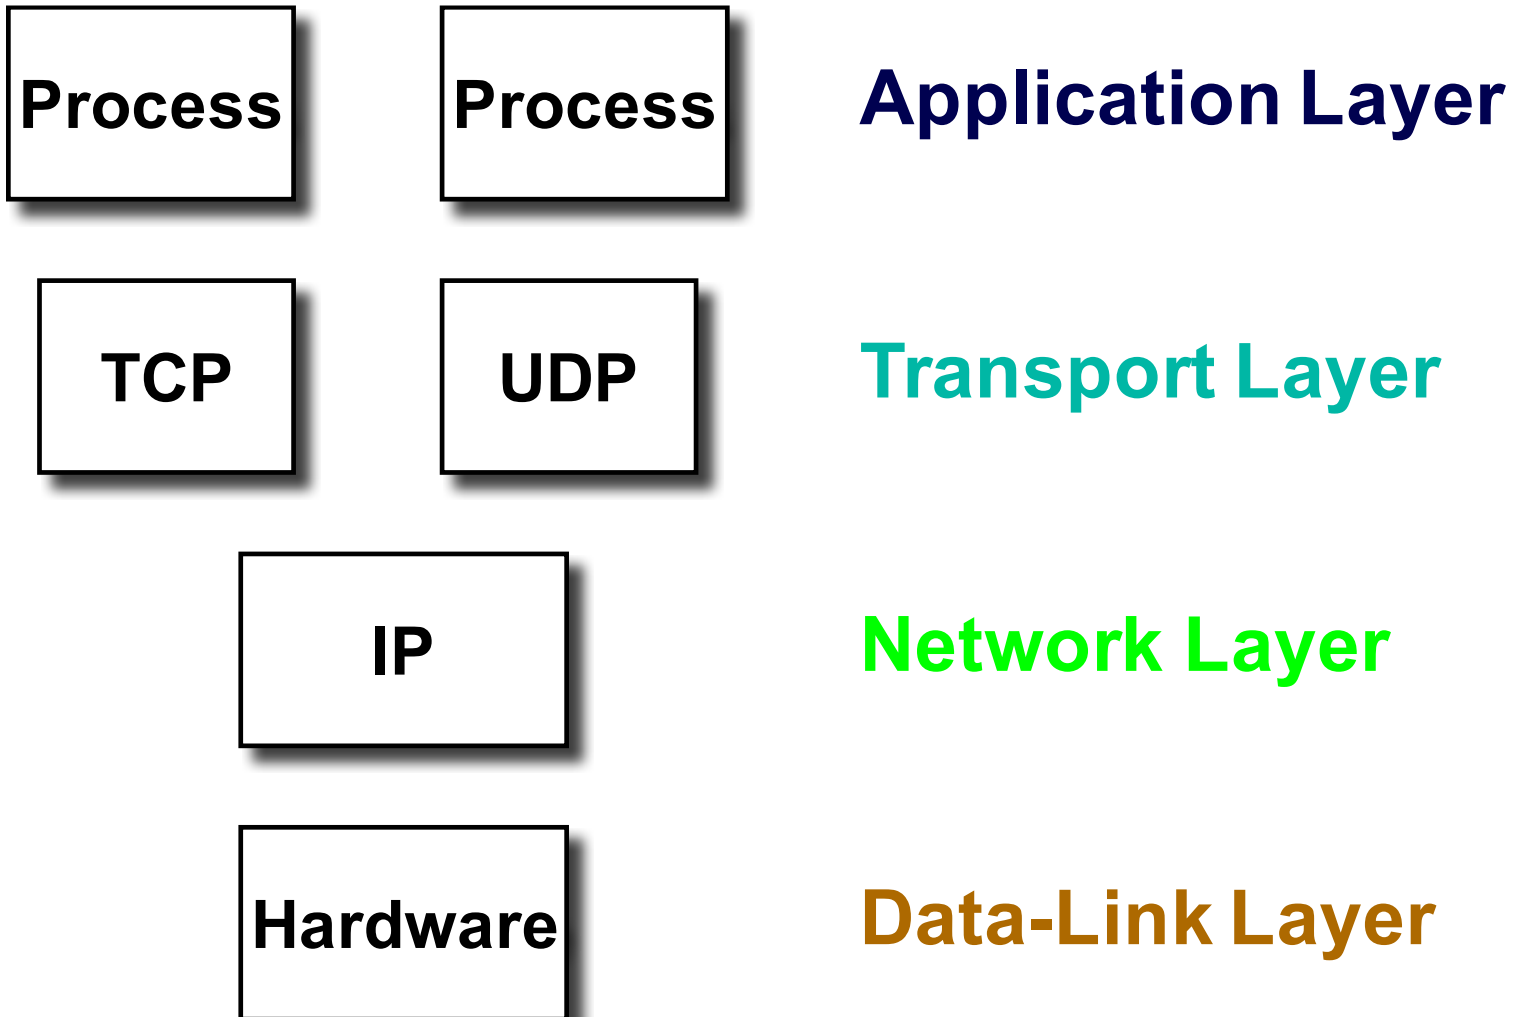
<https://tools.ietf.org/html/rfc791#section-3.1>



Datagram Fragmentation

- Each fragment has same structure
- IP requires reassembly done at destination only, not at intermediate routers
- Any lost fragments require ICMP error message be sent and entire datagram discarded

TCP / UDP over IP



UDP

- UDP is a transport protocol
 - Communication between two processes
- UDP uses IP to deliver datagrams to the proper host
- Uses *ports* to provide additional specification

UDP Format

<https://tools.ietf.org/html/rfc768>

Source Port	Destination Port
Length	Checksum
Data	

TCP

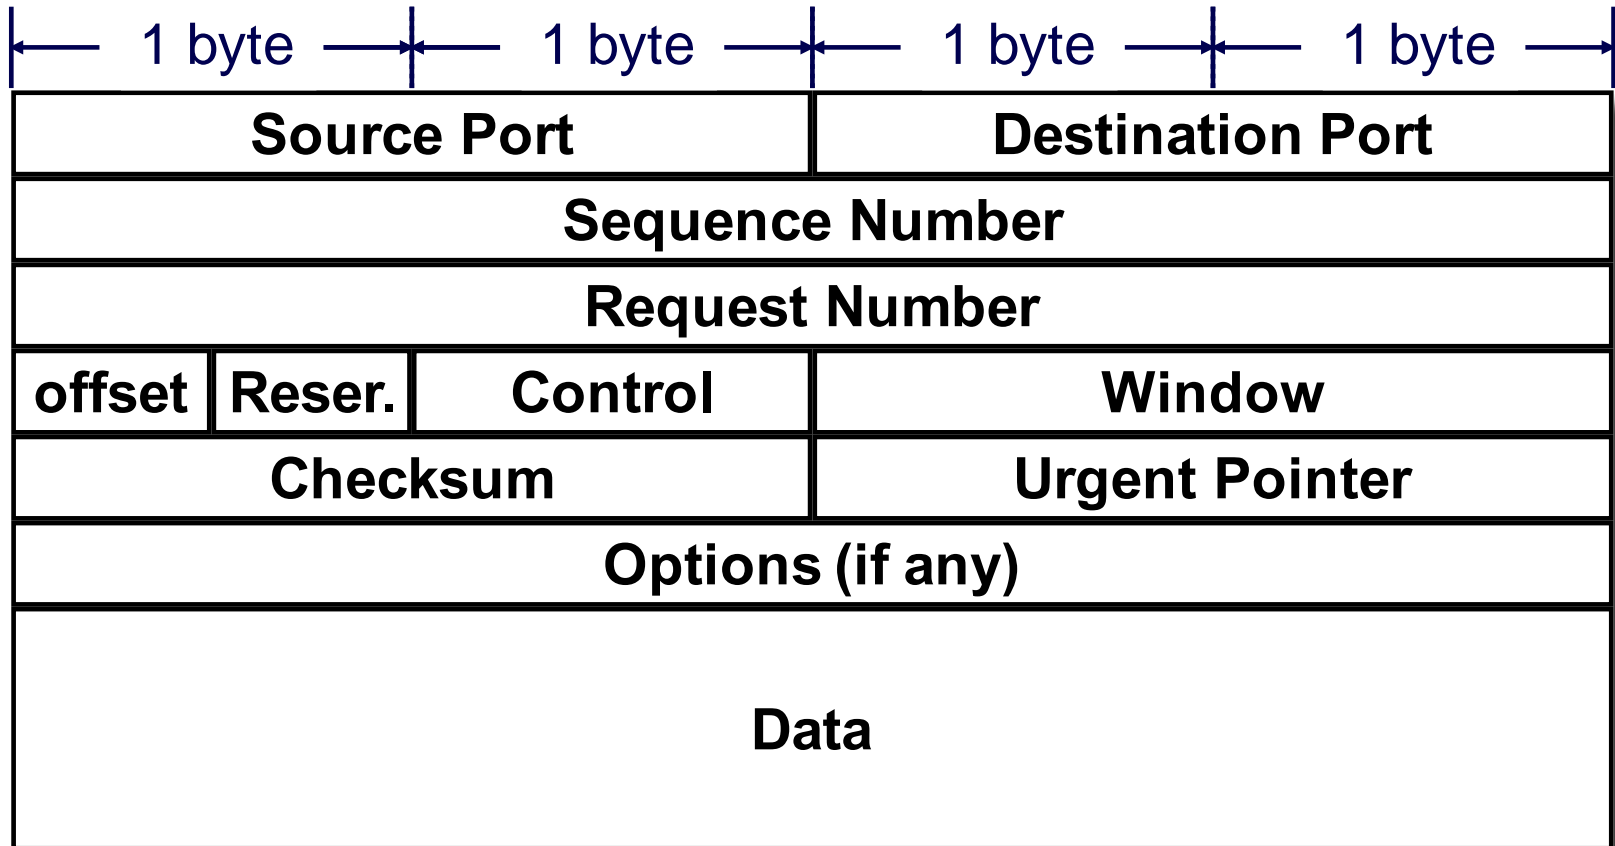
- TCP is a transport protocol
- In addition to all things provided by UDP, TCP provides:
 - Reliability
 - Full-duplex
 - Connection-oriented
 - Byte Stream

TCP Segments

- The chunk of data that TCP requests IP to transmit is called a *Segment*
- Each segment contains:
 - Data bytes from byte stream
 - Control information identifying data bytes

TCP Segment Format

<https://tools.ietf.org/html/rfc793#section-3.1>



UDP Sockets

Network Programming

Overview

- UDP (User Datagram Protocol) is unreliable
 - Potentially delivered out of order (or not at all!)
 - Connectionless
- TCP (Transmission Control Protocol) provides reliability
 - In-order
 - Connection-oriented

UDP

- UDP is a transport protocol
 - Communication between two processes
- UDP uses IP to deliver datagrams to the proper host
- Uses *ports* to provide additional specification



UDP Format

<https://tools.ietf.org/html/rfc768>

Source Port	Destination Port
Length	Checksum
Data	

But first... Endianness

- Some machines went big-endian (68K), others little-endian (x86)



- From Gulliver's Travels

Network Issues

- When hosts exchange single-byte data types, no problem
 - What about 32-bit word?
 - Are these the same value?
 - 0x00000001 and 0x01000000
 - Problem!
- Values sent from big-endian machine would be interpreted incorrectly on the little-endian machine!

Network Byte Order

- Network defines big-endian to be the byte order
 - May be different from *host byte order*
- Translation *always required*
 - Even on big-endian machines
 - How do you know what type of machine your code may be compiled on in the future?

Byte Order Functions

- `#include <netinet/in.h>`
- `uint16_t htons(uint16_t hs);`
- `uint32_t htonl(uint32_t hl);`
- `uint16_t ntohs(uint16_t ns);`
- `uint32_t ntohl(uint32_t nl);`

Sockets

- Berkeley sockets implementation, originally from 4.2BSD (1983!)
 - Effectively became POSIX sockets
- Building blocks for modern network-enabled programs
- Simple API

socket

- `#include <sys/socket.h>`
- `int socket(int domain, int type, int protocol);`
- Just creates an endpoint, nothing more!
- domain typically `PF_INET / AF_INET`
- type: `SOCK_[STREAM,DGRAM,RAW]`
- protocol: just use 0 for system default for given domain / type

bind

- `int bind(int fd, struct sockaddr *addr, socklen_t len) ;`
- `fd`: must be returned by `socket()`
- `sa`: `sockaddr` containing IP / port
- `len`: length of passed-in `sockaddr`
- Servers call `bind` upon startup

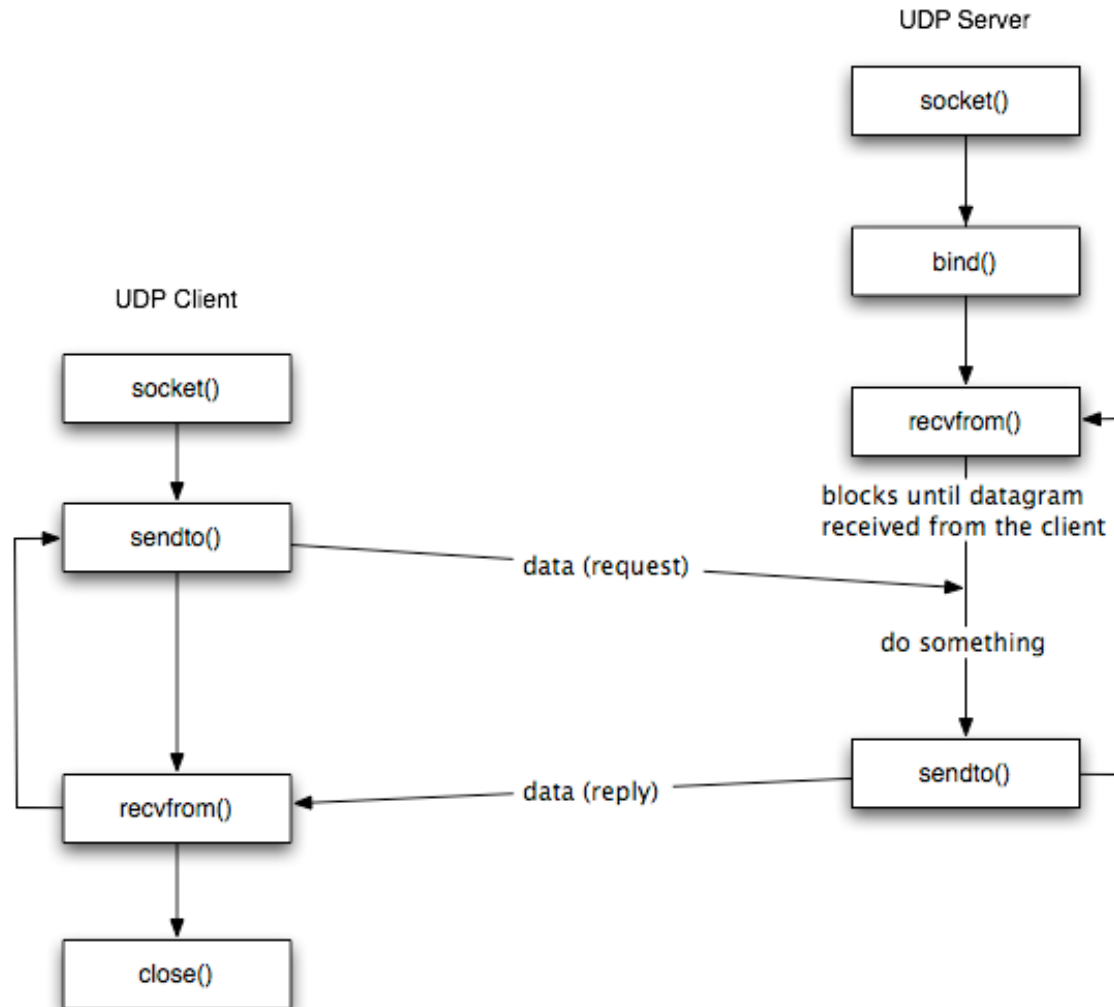
struct sockaddr_in

- struct sockaddr_in {
 __uint8_t sin_len;
 sa_family_t sin_family;
 in_port_t sin_port;
 struct in_addr sin_addr;
 char sin_zero[8];
};
- struct in_addr {
 in_addr_t s_addr;
};

Typical usage

```
struct sockaddr_in saddr;  
  
/* Zero out the memory */  
bzero(&saddr, sizeof(saddr));  
  
saddr.sin_family = PF_INET;  
saddr.sin_port = htons(1234);  
saddr.sin_addr.s_addr =  
    htonl(INADDR_ANY);
```

Typical UDP client/server



recvfrom / sendto

```
#include <sys/socket.h>
```

```
ssize_t recvfrom(int fd, void  
    *buf, size_t nbytes, int flags,  
    struct sockaddr *from,  
    socklen_t *len);
```

```
ssize_t sendto(int fd, void *buf,  
    size_t nbytes, int flags,  
    struct sockaddr *to, socklen_t len);
```

Oddities

- Sending 0 bytes is completely fine
 - 8 byte UDP header (no data)
- **recvfrom()** can return 0
 - Different from TCP where a 0 means peer has closed connection
- Both functions can also be used w/ TCP
 - But why???

udpserv01.c

```
#include "unp.h"

int
main(int argc, char **argv)
{
    int sockfd;
    struct sockaddr_in servaddr, cliaddr;

    sockfd = Socket(AF_INET, SOCK_DGRAM, 0);

    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family = AF_INET;
    servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
    servaddr.sin_port = htons(SERV_PORT);

    Bind(sockfd, (SA *) &servaddr, sizeof(servaddr));

    dg_echo(sockfd, (SA *) &cliaddr, sizeof(cliaddr));
}
```

dg_echo.c

```
#include "unp.h"

void
dg_echo(int sockfd, SA *pcliaddr, socklen_t clilen)
{
    int          n;
    socklen_t    len;
    char         mesg[MAXLINE];

    for ( ; ; ) {
        len = clilen;
        n = Recvfrom(sockfd, mesg, MAXLINE, 0, pcliaddr, &len);

        Sendto(sockfd, mesg, n, 0, pcliaddr, len);
    }
}
```

udpccli01.c

```
#include    "unp.h"

int
main(int argc, char **argv)
{
    int                sockfd;
    struct sockaddr_in servaddr;

    if (argc != 2)
        err_quit("usage: udpccli <IPaddress>");

    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family = AF_INET;
    servaddr.sin_port = htons(SERV_PORT);
    Inet_pton(AF_INET, argv[1], &servaddr.sin_addr);

    sockfd = Socket(AF_INET, SOCK_DGRAM, 0);

    dg_cli(stdin, sockfd, (SA *) &servaddr, sizeof(servaddr));

    exit(0);
}
```

dg_cli.c

```
#include "unp.h"

void
dg_cli(FILE *fp, int sockfd, const SA *pservaddr, socklen_t servlen)
{
    int n;
    char sendline[MAXLINE], recvline[MAXLINE + 1];
    while (Fgets(sendline, MAXLINE, fp) != NULL) {
        Sendto(sockfd, sendline, strlen(sendline), 0, pservaddr, servlen);

        n = Recvfrom(sockfd, recvline, MAXLINE, 0, NULL, NULL);

        recvline[n] = 0;    /* null terminate */
        Fputs(recvline, stdout);
    }
}
```

Not reliable in any way!

- What if messages get lost?
 - Client datagram?
 - Server datagram?
- How do we add reliability?
 - Timeouts (and retransmissions)
 - Sequence numbers
 - We'll have an assignment about this later in the semester

UDP connect ()

- We can call **connect ()** on a UDP socket
 - No longer able to use **sendto ()** but rather **write () / send ()**
 - Similarly for **recvfrom ()**; replace with **recv () / recvmsg ()**
- May improve performance if communicating with same host repeatedly

UDP Clients: DNS

- Quicker, no three-way handshake required (we'll revisit this later)
- Connectionless means less burden on the nameservers (we may revisit DNS later)
- Loss isn't terrible, just send another query some time later

DNS Resource Records

- A: hostname -> IPv4 address
- AAAA: hostname -> IPv6 address
- PTR: IP address -> hostname
- MX: Mail Exchange
- CNAME: canonical name

Skipping: Zeroconf/Bonjour

- Builds on DNS (and other tech as well)
- SRV records added to DNS
 - <https://tools.ietf.org/html/rfc2782>
- Designed to simplify networking
- Skipping for now because configuring issues (avahi)