

# Unit1

## C Programming and Intro to Linux

A crash course for  
C++/Windows Users

For your reference, not going to spend a lot of time here

## **C BASICS**

# C vs. C++: Differences

- C does not have classes/objects!
  - all code is in functions (subroutines).
- C structures can not have methods
- C I/O is based on library functions:
  - printf, scanf, fopen, fclose, fread, fwrite, ...

# C vs. C++: Differences (cont.)

- C does not support any function overloading (you can't have 2 functions with the same name).
- C does not have `new` or `delete`, you use `malloc()` and `free()` library functions to handle dynamic memory allocation/deallocation.
- C does not have *reference variables*

# Evolution of C

- Traditional C: 1978 by K&R
- Standard C: 1989 (aka ANSI C)
- Standard C: 1995, amendments to C89 standard
- Standard C: 1999, is the new definitive C standard replace all the others.
- GCC is a C99 compliant compiler (mostly, I think :-)).

# Standard C (C89)

- The addition of truly standard library
  - `libc.a`, `libm.a`, etc.
- New processor commands and features
- Function prototypes -- argument types specified in the function declaration
- New keywords: `const`, `volatile`, `signed`
- Wide chars, wide strings and multibyte characters.
- Clarifications to conversion rules, declarations and type checking

# Standard C (C95)

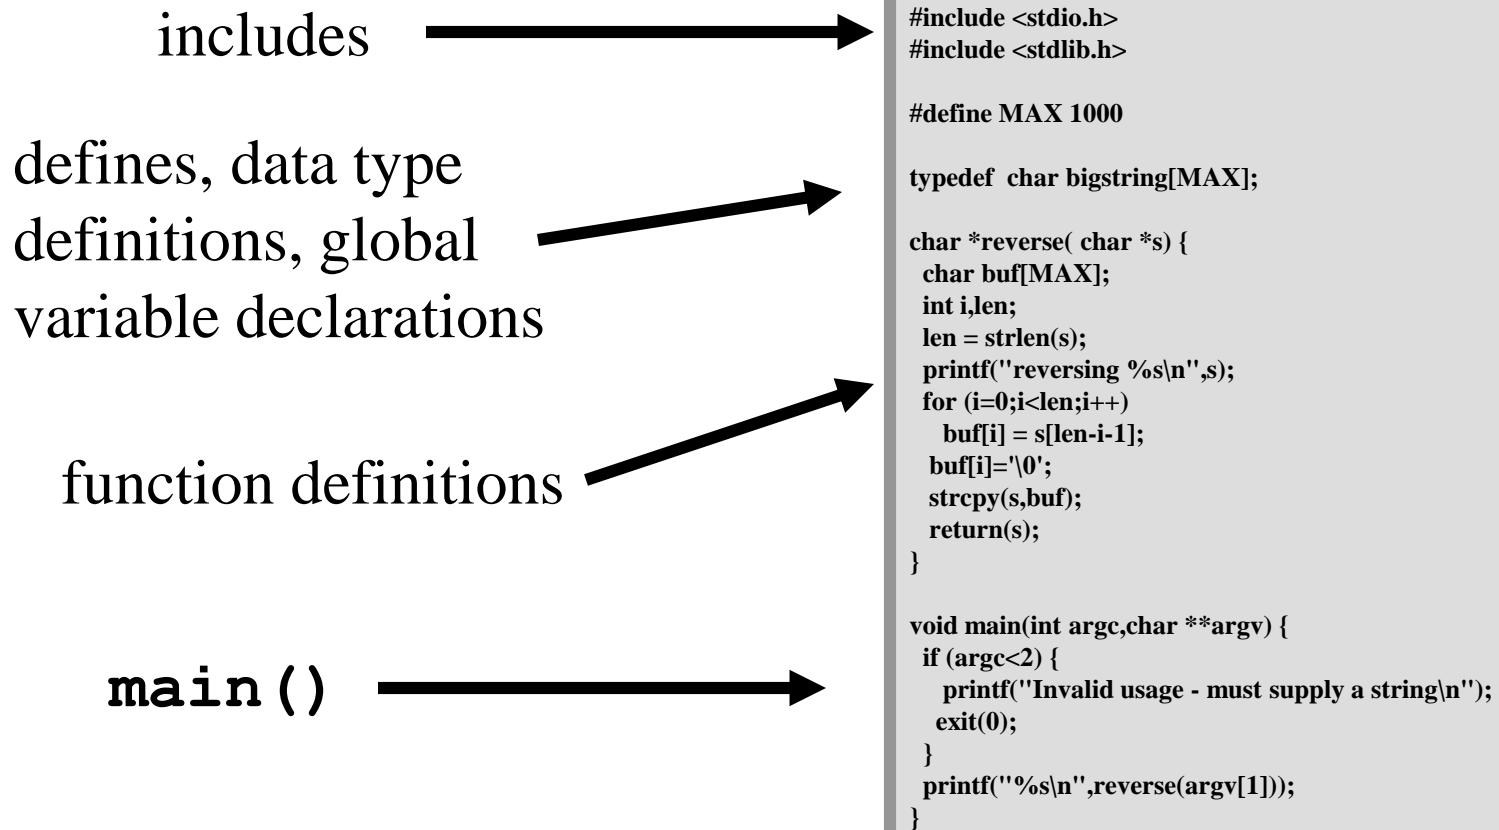
- 3 new std lib headers: `iso646.h`, `wctype.h` and `wchar.h`
- New formatting codes for `printf/scanf`
- A large number of new functions.

# Standard C (C99)

- Complex arithmetic
- Extensions to integer types, including the longer standard type (long long, long double)
- Boolean type (stdbool.h)
- Improved support for floating-point types, including math functions for all types
- C++ style comments (//)



# Typical C Program



# Using dynamic allocation

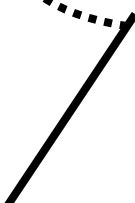
```
char *reverse( char *s) {  
    char *buf;  
    int i,len;  
    len = strlen(s);  
  
    /* allocate memory len + 1 for null term */  
    buf = (char *)malloc(len+1);  
    for (i=0;i<len;i++)  
        buf[i] = s[len-i-1];  
    buf[i]='\0';  
    strcpy(s,buf);  
    free(buf);  
    return(s);  
}
```

# Compiling on Unix

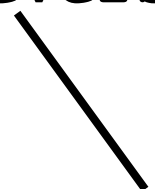
Traditionally the name of the C compiler that comes with Unix is “cc”.

We can use the Gnu compiler named “gcc”.

```
gcc -Wall -o reverse reverse.c
```



tells the compiler to  
create executable file  
with the name **reverse**



tells the compiler  
the name of the input  
file.

# Running the program

```
>./reverse Hello  
olleH
```

```
>./reverse This is a long string  
sihT
```

```
>./reverse "This is a long string"  
gnirts gnol a si sihT
```

# C Libraries

- Standard I/O: printf, scanf, fopen, fread, ...
- String functions: strcpy, strspn, strtok, ...
- Math: sin, cos, sqrt, exp, abs, pow, log,...
- System Calls: fork, exec, signal, kill, ...

# Quick I/O Primer - `printf`

```
int printf( const char *, . . . );
```

... means “variable number of arguments”.

The first argument is required (a string).

Given a simple string, `printf` just prints the string  
(to *standard output*).

# Simple printf

```
printf("Hi Dr. J. - I am a string\n");
```

```
printf("I\t have\t tabs\n");
```

```
char s[100];
```

```
strcpy(s, "printf is fun!\n");
```

```
printf(s);
```

# arguing with printf

You can tell `printf` to embed some values in the string – these values are determined at run-time.

```
printf("here is an integer: %d\n", i);
```

the `%d` is replaced by the value of the argument following the string (in this case `i`).



# More integer arguments

```
printf("%d + %d = %d\n",x,y,x+y) ;
```

```
for (j=99;j>=0;j--)
```

```
    printf("%d bottles of beer on the wall\n",  
j) ;
```

```
printf("%d is my favorite number\n",17) ;
```

# printf is dumb

- `%d` is replaced by the value of the parameter when treated as a integer.
- If you give `printf` something that is not an integer – it doesn't know!

```
printf("Print an int %d\n", "Hi Dr. J.");
```

```
Print an int 134513884
```

# Other formats

- `%d` is a format – it means “treat the parameter as a signed integer”
- `%u` means unsigned integer
- `%x` means print as hexadecimal
- `%s` means “treat it as a string”
- `%c` is for characters (`char`)
- `%f` is for floating point numbers
- `%%` means print a single ‘%’

# Special Format Characters

- `printf`
  - `%d`: signed integer
  - `%lld`: signed long long integer (64 bits)
  - `%u`: unsigned integer
  - `%x`: integer in hexadecimal format
  - `%f`: double
  - `%Lf`: long double
  - `%s` – a string
- `scanf` is the same except that `%f` is for float and `%Lf` is a double
- clang will “help” you if you use the wrong format

# Fun with printf

```
char *s = "Hi Dr. J";  
printf("The string \"%s\" is %d  
characters long\n",  
      s, strlen(s));  
  
printf("The square root of 10 is  
%f\n", sqrt(10));
```

# Controlling the output

- There are formatting options that you can use to control field width, precision, etc.

```
printf("The square root of 10  
is %20.15f\n", sqrt(10));
```

```
The square root of 10 is  
3.162277660168380
```

# Alas, we must move on

- There are more formats and format options for `printf`.
- The *man page* for `printf` includes a complete description (any decent C book will also).
- NOTE: to view the man page for `printf` you should type the following: `man 3 printf`

# Input - scanf

- `scanf` provides input from *standard input*.
- `scanf` is every bit as fun as `printf`!
- `scanf` is a little scary, you need to use pointers
- Actually, you don't really need pointers, just addresses.



# Remember Memory?

- Every C variable is stored in memory.
- Every memory location has an *address*.
- In C you can use variables called *pointers* to refer to variables by their address in memory.

# scanf

```
int scanf(const char *format, ...);
```

- Remember “...” means “variable number of arguments”
- Looks kinda like `printf`

# What scanf does

- Uses format string to determine what kind of variable(s) it should read.
- The arguments are the *addresses* of the variables.
- The & operator here means “Take the address of”:

```
int x, y;  
scanf ("%d %d", &x, &y) ;
```

# A simple example of `scanf`

```
float x;
```

```
printf("Enter a number\n");
```

```
scanf("%f", &x);
```

```
printf("The square root of %f is  
%f\n", x, sqrt(x));
```

# scanf and strings

Using `%s` in a `scanf` string tells `scanf` to read the next *word* from input – NOT a line of input:

```
char s[100]; // ALLOC SPACE!!  
printf("Type in your name\n");  
scanf("%s",s); // note: s is a char *  
printf("Your name is %s\n",s);
```

# Reading a line

- You can use the function `fgets` to read an entire line:

```
char *fgets(char *s, int size,  
            FILE *stream);
```

**size** is the maximum # chars

**FILE** is a *file handle*

# Using `fgets` to read from `stdin`

```
char s[101];
```

```
printf("Type in your name\n");  
fgets(s,100,stdin);
```

```
printf("Your name is %s\n",s);
```

# Other I/O stuff

- `fopen, fclose`
- `fscanf, fprintf, fgets`
- `fread, fwrite`
- Check the man pages for the details.



# String functions

```
char *strcpy(char *dest,  
              const char *src);
```

```
size_t strlen(const char *s);
```

```
char *strtok(char *s,  
              const char *delim);
```

# Math library

- The math library is often provided as an external library (not as part of the *standard C library*).
- You must tell the compiler you want the math library:

```
gcc -o myprog myprog.c -lm
```



means “add in the math library”

# Useful Predefined MACROS

- `__LINE__` : line # in source code file (%d)
- `__FILE__` : name of current source code file (%s)
- `__DATE__` : date “Mmm dd yyy” (%s)
- `__TIME__` : time of day, “hh:mm:ss” (%s)
- `__STDC__` : 1 if compiler is ISO compliant
- `__STDC_VERSION__` : integer (%s)

# O-O still possible with C

- Keyword **struct** is used to declare a record or “methodless” class, like

**struct Student**

```
{  
    char first_name[32];  
    char last_name[32];  
    unsigned int id;  
    double gpa;  
};
```

# Using a struct

```
int main()
{
    struct Student Suzy;
    /* init data members by hand */
    /* strcpy is a standard libC function */
    strcpy( Suzy.last_name, "Chapstick");
    Suzy.id = 12345;
}
```

# Variable Declarations

```
int main()
{
    struct Student Suzy;
    strcpy( Suzy.last_name, "Chapstick");
    int i; /* WRONG!! */
    struct Student Sam; /* WRONG !*/
    Suszy.id = 12345;
}
```

**All vars must be declared before the first executable statment in a function or block.**

**This has a significant impact on your “for” loops!**

# “for” Loops in C

- In C++ you will typically do:
  - `for( int i=0; i < num; i++ )`
- In C you MUST do:
  - `int i; /* top of scope */`
  - `for( i = 0; i < num; i++ )`
  - note, “i” exists outside of the for loop scope!
  - NO LONGER TRUE IN C99!!!

# Memory Allocation

- Use C system calls:
  - **void \*malloc(size\_t size)** returns a pointer to a chunk of memory which is the size in bytes requested
  - **void \*calloc(size\_t nmemb, size\_t size)** same as malloc but puts zeros in all bytes and asks for the number of elements and size of an element.
  - **void free(void \*ptr)** deallocates a chunk of memory. Acts much like the delete operator.
  - **void \*realloc(void \*ptr, size\_t size)** changes the size of the memory chunk point to by **ptr** to **size** bytes.
  - prototypes found in the **stdlib.h** header file.



# Example Memory Allocation

```
void main()
{
    char *cptr;
    double *dblptr;
    struct Student *stuptr;
    /* equiv to cptr = new char[100]; */
    cptr = (char *) malloc(100);
    /* equiv to dblptr = new double[100]; */
    dblptr = (double *) malloc(sizeof(double) * 100);
    /* equiv to stuptr = new Student[100]; */
    stuptr = (struct Student *) malloc( sizeof( struct Student) *
    100);
}
```

# Well Used Header Files (based on Linux)

- ***stdio.h*** – printf/scanf/ FILE type
- ***stdlib.h*** – convert routines like string-to-XX, (strtol, strtod, stro), rand num gen, calloc and malloc
- ***unistd.h*** – system calls like fork, exec, read, write
- ***math.h/float.h*** – math routines
- ***errno.h*** – standard error numbers for return values and error handling routines like perror, system call numbers (in Linux).

This section is for your reference, not going to spend any time here

## **UNIX BASICS**

# Unix Accounts

- To access a Unix system you need to have an *account*.
- Unix account includes:
  - username and password
  - userid and groupid
  - home directory
  - shell

# Home Directory

- A home directory is a place in the file system where files related to an account are stored.
- A *directory* is like a Windows folder (more on this later).
- Many unix commands and applications make use of the account home directory (as a place to look for customization files).

# Shell

- A Shell is a unix program that provides an interactive session - a text-based user interface.
- When you log in to a Unix system, the program you initially interact with is your shell.
- There are a number of popular shells that are available.

# Your Home Directory

- Every Unix process\* has a notion of the “current working directory”.
- Your shell (which is a process) starts with the current working directory set to your home directory.

\*A process is an instance of a *program* that is currently running.

# Interacting with the Shell

- The shell prints a prompt and waits for you to type in a command.
- The shell can deal with a couple of types of commands:
  - shell internals - commands that the shell handles directly.
  - External programs - the shell runs a program for you.



# Files and File Names

- A file is a basic unit of storage (usually storage on a disk).
- Every file has a name.
- Unix file names can contain any characters\* (although some make it difficult to access the file).
- Unix file names can be long!
  - how long depends on your specific flavor of Unix

\*except NUL and / but your filesystem may have additional restrictions

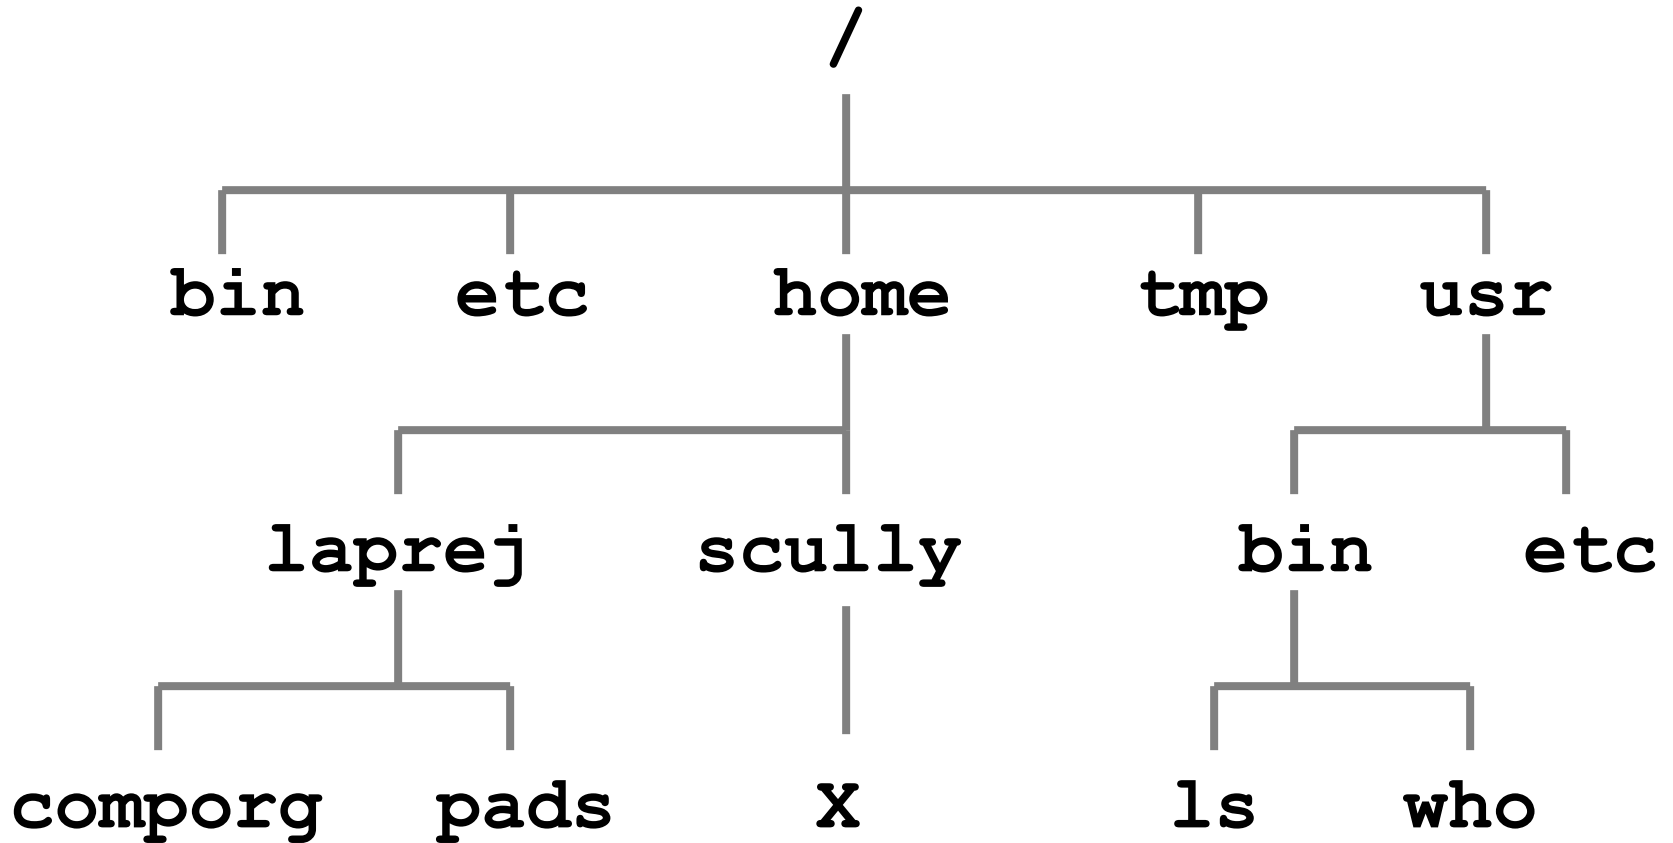
# File Contents

- Each file can hold some raw data.
- Unix does not impose any structure on files
  - files can hold any sequence of bytes.
- Many programs *interpret* the contents of a file as having some special structure
  - text file, sequence of integers, database records, etc.

# Directories

- A directory is a special kind of file - Unix uses a directory to hold information about other files.
- We often think of a directory as a container that holds other files (or directories).
- Mac and Windows users: A directory is the same idea as a *folder*.
- *Folders are used as a GUI interface to directories and not unique to Unix/Linux/FreeBSD*

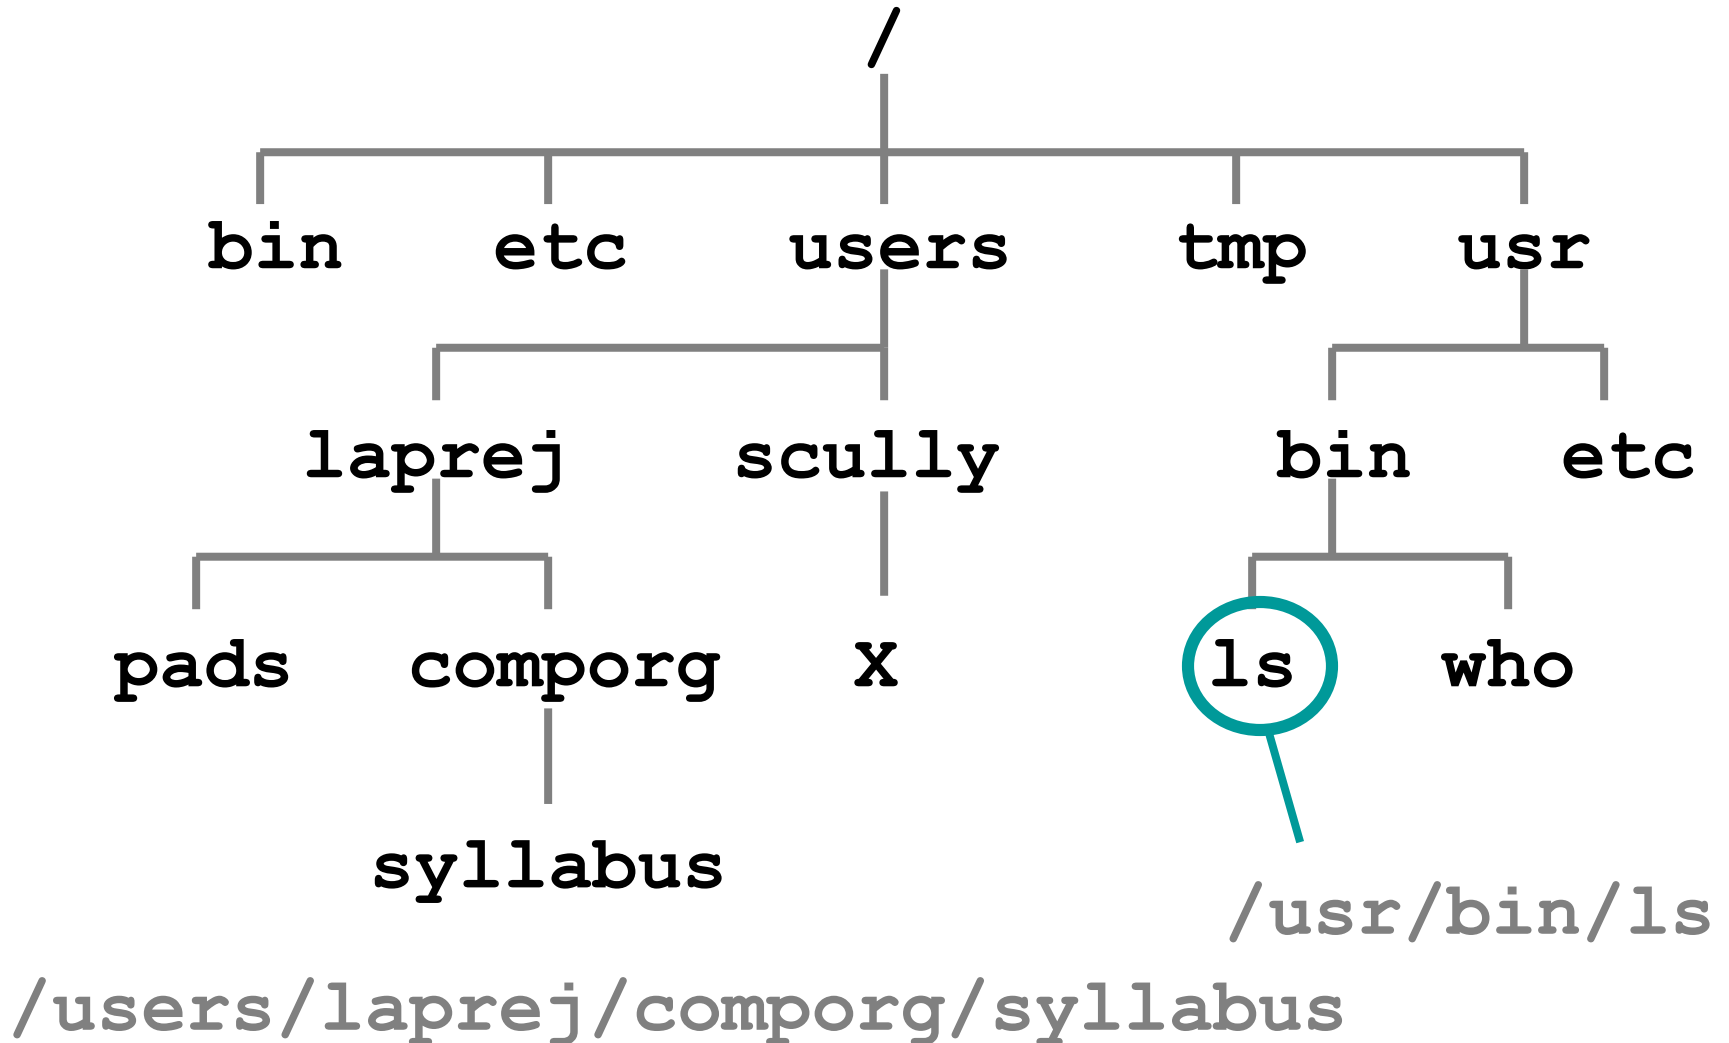
# The Filesystem



# Unix Filesystem

- The filesystem is a hierarchical system of organizing files and directories.
- The top level in the hierarchy is called the "root" and *holds* all files and directories.
- The name of the root directory is `/`

# Pathname Examples



# Absolute Pathnames

- The pathnames described in the previous slides start at the *root*.
- These pathnames are called "absolute pathnames".
- We can also talk about the pathname of a file *relative* to a directory.

# Relative Pathnames

- If we are *in* the directory `/home/laprej`, the relative pathname of the file **`syllabus`** in the directory `/home/laprej/comporg/` is:

**`comporg/syllabus`**

- Most Unix commands deal with pathnames!
- We will usually use relative pathnames when specifying files.



# The current directory and *parent* directory

- There is a special relative pathname for the current directory:

.

- There is a special relative pathname for the parent directory:

..

# Your home directory

- There is also a special relative pathname for the current user's home directory:

~

- Try this:

```
touch /home/yourusername/afile
```

```
ls -l /home/yourusername
```

```
touch -l ~/anotherfile
```

```
ls ~
```

# The `ls` command

- The `ls` command displays the names of some files.
- If you give it the name of a directory as a *command line parameter* it will list all the files in the named directory.

# ls Command Line Options

- We can modify the output format of the **ls** program with a *command line option*.
- The **ls** command support a bunch of options:
  - **l** *long* format (include file times, owner and permissions)
  - **a** *all* (shows hidden\* files as well as regular files)
  - **F** include special char to indicate file types.

\*hidden files have names that start with "."

# Moving Around in the Filesystem

- The `cd` command can change the current working directory:

`cd`            *change directory*

- The general form is:

`cd [directoryname]`

# cd

- With no parameter, the **cd** command changes the current directory to your home directory.
- You can also give **cd** a relative or absolute pathname:

```
cd /usr
```

```
cd ..
```

# Copying Files

- The **cp** command copies files:

**cp [options] source dest**

- The source is the name of the file you want to copy.
- dest is the name of the new file.
- source and dest can be relative or absolute.
- -r is a common option for "recursive", needed if the source is a directory

## Another form of **cp**

- If you specify a dest that is a directory, cp will put a copy of the source in the directory.
- The filename will be the same as the filename of the source file.

```
cp [options] source destdir
```



# Deleting (removing) Files

- The **rm** command deletes files:

```
rm [options] names...
```

- **rm** stands for "remove".
- You can remove many files at once:

```
rm foo /tmp/blah /users/clinton/intern
```

# File Permissions

- Each file has a set of permissions that control who can mess with the file.
- There are three kinds of permissions:
  - read                      abbreviated **r**
  - write                      abbreviated **w**
  - execute                      abbreviated **x**
- There are separate permissions for the file owner, group owner and everyone else.

# ls -l and permissions

**-** **rwx** **rwx** **rwx**

Owner

Group

Others

**Type of file:**

**-** means plain file

**d** means directory

# ***rwX***

- Files:
  - **r**: allowed to read.
  - **w**: allowed to write.
  - **x**: allowed to execute
- Directories:
  - **r**: allowed to see the names of the files.
  - **w**: allowed to add and remove files.
  - **x**: allowed to enter the directory

# Changing Permissions

- The **chmod** command changes the permissions associated with a file or directory.
- There are a number of forms of chmod, this is the simplest:

**chmod mode file**

# chmod mode file

- Mode has the following form\*:

**[u go a] [+ - =] [rwx]**

u=user      g=group      o=other      a=all

+ add permission      - remove permission      = set permission

\*The form is really more complicated, but this simple version will do enough for now.

# chmod examples

```
> ls -al foo  
rwxrwx--x    1 laprej grads ...
```

```
> chmod g-x foo
```

```
> ls -al foo  
-rwxrw---x    1 laprej grads
```

```
> chmod u-r .
```

```
> ls -al foo
```

```
ls: .: Permission denied
```

## Other filesystem and file commands

- **mkdir** make directory
- **rmdir** remove directory
- **touch** change file timestamp (can also create a blank file)
- **cat** concatenate files and print out to terminal.



# Shells

Also known as: Unix Command  
Interpreter

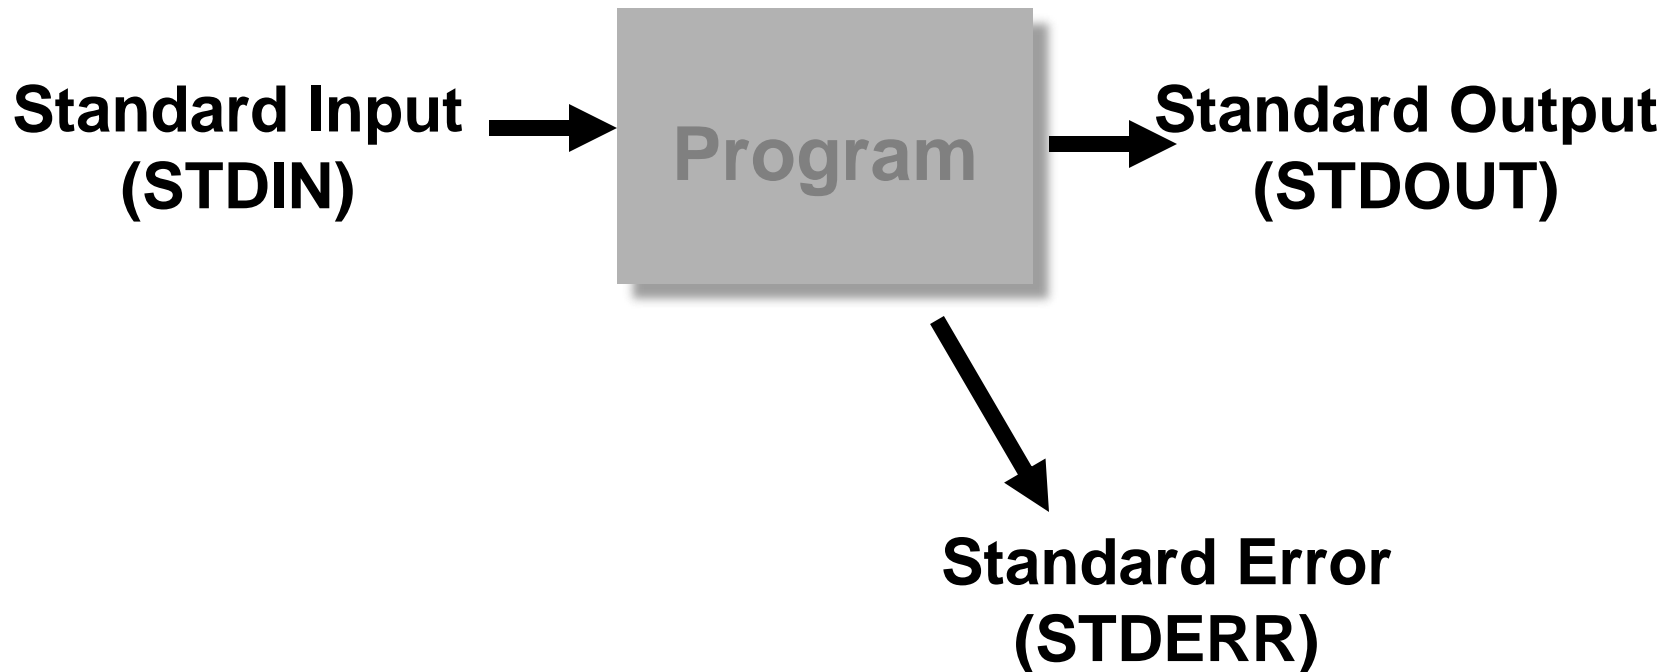
# Shell as a user interface

- A shell is a command interpreter that turns text that you type (at the command line) into actions:
  - runs a program, perhaps the `ls` program.
  - allows you to edit a *command line*.
  - can establish alternative sources of input and destinations for output for programs.

# Running a Program

- You type in the name of a program and some command line options:
  - The shell reads this line, finds the program and runs it, feeding it the options you specified.
  - The shell establishes 3 I/O *channels*:
    - Standard Input
    - Standard Output
    - Standard Error

# Programs and Standard I/O



# Unix Commands

- Most Unix commands (programs):
  - read something from standard input.
  - send something to standard output (typically depends on what the input is!).
  - send error messages to standard error.

# Defaults for I/O

- When a shell runs a program for you:
  - standard input is your keyboard.
  - standard output is your screen/window.
  - standard error is your screen/window.

# Terminating Standard Input

- If standard input is your keyboard, you can type stuff in that goes to a program.
- To end the input you press Ctrl-D (^D) on a line by itself, this ends the input *stream*.
- The shell is a program that reads from standard input.
- What happens when you give the shell ^D?

# The special character \*

- \* matches anything.
- If you give the shell \* by itself (as a command line argument) the shell will remove the \* and replace it with all the filenames in the current directory.
- “**a\*b**” matches all files in the current directory that start with **a** and end with **b**.



# \* and ls

- Things to try:

```
ls *
```

```
ls -al *
```

```
ls a*
```

```
ls *b
```

# Input Redirection

- The shell can attach things other than your keyboard to standard input.
  - A file (the contents of the file are fed to a program as if you typed it).
  - A pipe (the output of another program is fed as input as if you typed it).

# Output Redirection

- The shell can attach things other than your screen to standard output (or stderr).
  - A file (the output of a program is stored in file).
  - A pipe (the output of a program is fed as input to another program).

# How to tell the shell to redirect things

- To tell the shell to store the output of your program in a file, follow the command line for the program with the “>” character followed by the filename:

```
ls > lsout
```

the command above will create a file named **lsout** and put the output of the **ls** command in the file.

# Input redirection

- To tell the shell to get standard input from a file, use the “<” character:

**sort < nums**

- The command above would sort the lines in the file `nums` and send the result to `stdout`.

# You can do both!

```
sort < nums > sortednums
```

```
tr a-z A-Z < letter > rudeletter
```

Note: “tr” command is translate.

Here it replaces all letters  
“a-z” with “A-Z”

# More Output redirection

- To tell the shell to print standard error to a file, use the “2>” phrase:

```
gcc buggy_file.c 2> compile.log
```

- The command above would send any error messages during the compile to compile.log


# Even More Output redirection

- To tell the shell to print standard error AND standard out to the same file, use the “&>” phrase:  
**`./kind_of_works.out &> run.log`**
- The command above would send any output and errors to run.log



# Pipes

- A pipe is a holder for a stream of data.
- A pipe can be used to hold the output of one program and feed it to the input of another.



# Asking for a pipe

- Separate 2 commands with the “|” character.
- The shell does all the work!

```
ls | sort
```

```
ls | sort > sortedls
```

# The **PATH**

- Each time you give the shell a command line it does the following:
  - Checks to see if the command is a shell built-in.
  - If not - tries to find a program whose name (the filename) is the same as the command.
- The **PATH** variable tells the shell where to look for programs (non built-in commands).

# echo \$PATH

```
===== [foo.cs.rpi.edu] - 22:43:17 =====  
/home/laprej/comporg echo $PATH  
/home/laprej/bin:/usr/bin:/bin:/usr/local/bin  
:/usr/sbin:/usr/bin/X11:/usr/games:/usr/local/packages/netscape
```

- The **PATH** is a list of ":" delimited directories.
- The **PATH** is a list and a *search order*.
- You can add stuff to your PATH by changing the shell startup file (`~/ .bashrc`)

# Job Control

- The shell allows you to manage *jobs*
  - place *jobs* in the *background*
  - move a job to the foreground
  - suspend a job
  - kill a job

# Background jobs

- If you follow a command line with "&", the shell will run the *job* in the background.
  - you don't need to wait for the job to complete, you can type in a new command right away.
  - you can have a bunch of jobs running at once.
  - you can do all this with a single terminal (window).

```
ls -lR > saved_ls &
```

# Listing jobs

- The command *jobs* will list all background jobs:

```
> jobs
```

```
[1] Running      ls -lR > saved_ls &
```

```
>
```

- The shell assigns a number to each job (this one is job number 1).

# Suspending and Killing the Foreground Job

- You can suspend the foreground job by pressing ^Z (Ctrl-Z).
  - Suspend means the job is stopped, but not dead.
  - The job will show up in the **jobs** output.
- You can *kill* the foreground job by pressing ^C (Ctrl-C).
- If ^C does not work, use ^Z to get back to your terminal prompt and issue:  
\$> **kill -9 %1**



# Quoting - the problem

- We've already seen that some characters mean something special when typed on the command line: `*` (also `?`, `[]`)
- What if we don't want the shell to treat these as special - we really mean `*`, not all the files in the current directory:

```
echo here is a star *
```

# Quoting - the solution

- To turn off special meaning - surround a string with double quotes:

➤ `echo here is a star "*"`

➤ `here is a star *`

# Quoting Exceptions

- Some *special* characters are **not** ignored even if inside double quotes:
- \$ (prefix for variable names)
- " the quote character itself
- \ slash is always something special (\n)
  - you can use \\$ to mean \$ or \" to mean "

```
>echo "This is a quote \" "  
>This is a "
```

# Single quotes

- You can use single quotes just like double quotes.
  - Nothing (except ' ) is treated special.

```
> echo 'This is a quote \" '
```

```
> This is a quote \"
```

# What are stdin, stdout, stderr?

- File descriptors...or more precisely a pointer to type FILE.
- These FILE descriptors are setup when your program is run.
- So, then what about regular user files...

# File I/O Operations

- `fopen` -- opens a file
- `fclose` -- close a file
- `fprintf` -- “printf” to a file.
- `fscanf` -- read input from a file.
- ...and many other routines..

# fopen

```
#include<stdio.h>
```

```
void main()
```

```
{
```

```
    FILE *myfile;
```

```
    myfile = fopen( "myfile.txt", "w");
```

```
}
```

- 2nd arg is mode:
  - w -- create/truncate file for writing
  - w+ -- create/truncate for writing and reading
  - r -- open for reading
  - r+ -- open for reading and writing

# fclose

```
#include<stdio.h>
#include<errno.h>
void main()
{
    FILE *myfile;
    if( NULL == (myfile = fopen( "myfile.txt", "w")) )
    {
        perror("fopen failed in main");
        exit(-1);
    }
    fclose( myfile );
    /* could check for error here, but usually not
    needed */
}
```



# fscanf

```
#include<stdio.h>
#include<errno.h>
void main()
{
    FILE *myfile;
    int i, j, k;
    char buffer[80];
    if( NULL == (myfile = fopen( "myfile.txt", "w")) )
    {
        perror("fopen failed in main");
        exit(-1);
    }
    fscanf( myfile, "%d %d %d %s", &i, &j, &k, buffer);
    fclose( myfile );
    /* could check for error here, but usually not needed */
}
```

# sscanf

```
#include<stdio.h>
#include<errno.h>
void main()
{
    FILE *myfile;
    int i, j, k;
    char buffer[1024];
    char name[80];
    if( NULL == (myfile = fopen( "myfile.txt", "w")) )
    {
        perror("fopen failed in main");
        exit(-1);
    }
    fgets( buffer, 1024, myfile );
    sscanf( buffer, "%d %d %d %s", &i, &j, &k, name);
    fclose( myfile );
    /* could check for error here, but usually not needed */
}
```

# fprintf

```
#include<stdio.h>
#include<errno.h>
void main()
{
    FILE *myfile;
    int i, j, k;
    char buffer[80];
    if( NULL == (myfile = fopen( "myfile.txt", "w")))
    {
        perror("fopen failed in main");
        exit(-1);
    }
    fscanf( myfile, "%d %d %d %s", &i, &j, &k, buffer);
    fprintf( myfile, "%d %d %d %s, i, j, k, buffer );
    fclose( myfile );
    /* could check for error here, but usually not needed */
}
```

# Pipes

- They to are realized as a file descriptor which links either ouput to input or input to output.
  - recall doing shell commands of the form:
  - `> ls -al | grep "Jan 1" | more`
  - “|” is implemented as a libc call to “popen”

# Operating Systems: Unix/Linux

---

# Posix - Portable Operating System Interface

- Posix is a popular standard for Unix-like operating systems.
- Posix is actually a *collection* of standards that cover system calls, libraries, applications and more...
- Posix 1003.1 defines the C language interface to a Unix-like kernel.

# Posix and Unix

- Most current Unix-like operating systems are *Posix compliant* (or nearly so).

Linux, BSD, Mac OS X

- We won't do anything fancy enough that we need to worry about specific versions/flavors of Unix (any Unix will do).

# Posix 1003.1

- process primitives
  - creating and managing processes
- managing process environment
  - user ids, groups, process ids, etc.
- file and directory I/O
- terminal I/O
- system databases (passwords, etc)



# System Calls

- A *system call* is an interface to the kernel that makes some request for a service.
- The actual implementation (how a program actually contacts the operating system) depends on the specific version of Unix and the processor.
- The C interface to system calls is standard (so we can write an program and it will work anywhere).

# opendir/closedir

```
#include <dirent.h>
```

```
DIR *opendir(const char *dirname);
```

```
int closedir(DIR *dirp);
```

**opendir()** opens a directory, just like **fopen()** but for directories

**closedir()** is like **fclose()**

# readdir

```
#include <dirent.h>  
struct dirent *readdir(DIR *dirp) ;
```

Gets the next entry in the directory,  
which can be "." or ".." so be  
careful!

# Directory example part 1

```
#include <stdio.h>
#include <dirent.h>
#include <unistd.h>

int main() {
    DIR* dirp;
    struct dirent *dp;
    char dirname[255];

    if((dirp = opendir(".")) == NULL) {
        printf("Couldn't open dir\n");
        return -1;
    }
```

# Directory example part 2

```
printf("The contents of %s are:\n",
        getcwd(dirname,255));
while (dp != NULL) {
    if((dp = readdir(dirp)) != NULL) {
        printf("\t%s\n",dp->d_name);
    }
}

closedir(dirp);
return 0;
}
```

# stat

```
#include <sys/stat.h>
```

```
int stat(const char *restrict path,  
struct stat *restrict buf);
```

**stat** structure documented in **man 2**

**stat** (if you just do **man stat** you will probably get the page in section 1 which is for the stat program)

# stat example

```
while (dp != NULL) {
    if((dp = readdir(dirp)) != NULL) {
        if(stat(dp->d_name, &sb) == -1) {
            printf("\t !!! Failed to stat %s\n", dp-
>d_name);
            continue;
        }
        if(sb.st_mode & S_IFDIR) {
            printf("\t D %s\n", dp->d_name);
        }
        else if(sb.st_mode & S_IFDIR) {
            printf("\t F %s\n", dp->d_name);
        }
        else{
            printf("\t ? %s\n", dp->d_name);
        }
    }
}
```

# stat example v2

```
while (dp != NULL) {
    if((dp = readdir(dirp)) != NULL) {
        if(stat(dp->d_name, &sb) == -1) {
            printf("\t !!! Failed to stat %s\n", dp-
>d_name);
            continue;
        }
        if(S_ISDIR(sb.st_mode)) {
            printf("\t D %s\n", dp->d_name);
        }
        else if(S_ISREG(sb.st_mode)) {
            printf("\t F %s\n", dp->d_name);
        }
        else{
            printf("\t ? %s\n", dp->d_name);
        }
    }
}
```



# Unix Processes

- Every process has the following attributes:
  - a *process id* (a small integer)
  - a *user id* (a small integer)
  - a *group id* (a small integer)
  - a *current working directory*.
  - a chunk of memory that hold name/value pairs as text strings (the *environment variables*).
  - a bunch of other things...

# Creating a Process

- The only way to create a new process is to issue the `fork()` system call.
- `fork()` *splits* the current process in to 2 processes, one is called the *parent* and the other is called the *child*.

# Parent and Child Processes

- The child process is a *copy* of the parent process.
- Same program.
- Same place in the program (almost – we'll see in a second).
- The child process gets a new process ID.

# Process Inheritance

- The child process *inherits* many attributes from the parent, including:
  - current working directory
  - user id
  - group id

# The `fork()` system call

```
#include <unistd.h>
pid_t fork(void);
```

`fork()` returns a process id (a small integer).

`fork()` returns twice!

In the parent – `fork` returns the id of the child process.

In the child – `fork` returns a 0.

# Example

```
#include <unistd.h>
#include <stdio.h>

void main(void) {
    if (fork())
        printf("I am the parent\n");
    else
        printf("I am the child\n");
    printf("I am the walrus\n");
}
```

# Bad Example (don't try this!)

```
#include <unistd.h>
#include <stdio.h>

void main(void) {
    while (fork()) {
        printf("I am the parent %d\n"
               ,getpid());
    }
    printf("I am the child %d\n"
           ,getpid());
}
```

*fork bomb!*

# I told you so...

- Try pressing Ctrl-C to stop the program.
- It might be too late.
- If this is your own machine – try rebooting.
- If this is a campus machine – run for your life.  
If they catch you – deny everything.



# Switching Programs

- `fork()` is the only way to create a new process.
- This would be almost useless if there was not a way to switch what *program* is associated with a process.
- The `exec()` system call is used to start a new program.

# exec

- There are actually a number of exec functions:  
`execlp execl execl_e execvp execv execve`
- The difference between functions is the parameters... (how the new program is identified and some attributes that should be set).

# The exec family

- When you call a member of the exec family you give it the pathname of the executable file that you want to run.
- If all goes well, exec will never return!
- The process *becomes* the new program.

# execl()

```
int execl(char *path,  
          char *arg0,  
          char *arg1, ...,  
          char *argn,  
          (char *) 0) ;  
  
execl("/home/laprej/reverse",  
      "reverse", "Hello!", NULL) ;
```

# A complete `exec1` example

```
#include <unistd.h>  /* exec, getcwd */
#include <stdio.h>    /* printf */

/* Exec example code */
/* This program simply execs "/bin/ls" */

void main(void) {
    char buf[1000];

    printf("Here are the files in %s:\n",
           getcwd(buf,1000));
    exec1("/bin/ls","ls","-al",NULL);
    printf("If exec works, this line won't be
    printed\n");
}
```

# `fork()` and `exec()` together

- Program does the following:
  - `fork()` - results in 2 processes
  - parent prints out it's **PID** and waits for child process to finish (to exit).
  - child prints out it's **PID** and then **execs** "**ls**" and exits.

# execandfork.c part 1

```
#include <unistd.h>    /* exec, getcwd */
#include <stdio.h>      /* printf */
#include <sys/types.h>  /* need for wait */
#include <sys/wait.h>   /* wait() */
```

# execandfork.c part 2

```
void child(void) {  
    int pid = getpid();  
  
    printf("Child process PID is %d\n",pid);  
    printf("Child now ready to exec ls\n");  
    execl("/bin/ls","ls",NULL);  
}
```



# execandfork.c part 3

```
void parent(void) {  
    int pid = getpid();  
    int stat;  
  
    printf("Parent process PID is %d\n",pid);  
    printf("Parent waiting for child\n");  
    wait(&stat);  
    printf("Child is done. Parent now  
    transporting to the surface\n");  
}
```

# execandfork.c part 4

```
void main(void) {  
    printf("In main - starting things with a  
fork()\n");  
    if (fork()) {  
        parent();  
    } else {  
        child();  
    }  
    printf("Done in main()\n");  
}
```

# execandfork.c output

```
> ./execandfork
```

```
In main - starting things with a fork()
```

```
Parent process PID is 759
```

```
Parent process is waiting for child
```

```
Child process PID is 760
```

```
Child now ready to exec ls
```

```
exec          execandfork      fork
```

```
exec.c        execandfork.c    fork.c
```

```
Child is done. Parent now transporting to  
the surface
```

```
Done in main()
```

```
>
```