

# Sockets Overview

---

Network Programming

# Overview

- UDP (User Datagram Protocol) is unreliable
  - Potentially delivered out of order (or not at all!)
  - Connectionless
- TCP (Transmission Control Protocol) provides reliability
  - In-order delivery of packets
  - Connection-oriented
- Both sit on top of another protocol

# Sockets

- Berkeley sockets implementation, originally from 4.2BSD (1983!)
  - Effectively became POSIX sockets
- Building blocks for modern network-enabled programs
- Simple API
  - Typical server / client sequence to follow...

# Server

- Acquire socket with `socket()` call
- Bind socket to address (or wildcard) with `bind()`
- Cause socket to enter listening state with `listen()`
- `accept()`: block until connection request

# socket

- `#include <sys/socket.h>`
- `int socket(int domain, int type, int protocol);`
- Just creates an endpoint, nothing more!
- domain typically `PF_INET / AF_INET`
- type: `SOCK_[STREAM,DGRAM,RAW]`
- protocol: just use 0 for system default for given domain / type

# bind

- `int bind(int fd, struct sockaddr *addr, socklen_t len) ;`
- `fd`: must be returned by `socket()`
- `sa`: `sockaddr` containing IP / port
- `len`: length of passed-in `sockaddr`
- Servers call `bind` upon startup

# listen

- `int listen(int fd, int p_log);`
- Cause fd to enter listening state (after call to bind)
- Connections aren't instantaneous. Allows for p\_log pending connections

# accept

- `int accept(int fd, struct sockaddr *addr, socklen_t *addr_len) ;`
- `fd`: accept incoming connections
- `addr`: filled in with remote host info
- `addr_len`: initially should contain amount of space in `addr`, on return will contain length of filled-in `addr` data



# Client

---

- Acquire socket with `socket()` call
- Call `connect()` to connect to the server

# connect

- `int connect(int fd, struct sockaddr *sa, socklen_t len);`
- `fd`: must be returned by `socket()`
- `sa`: `sockaddr` containing IP / port
- `len`: length of passed-in `sockaddr`
- Clients call `connect()` and not `bind()`
- TCP: three way handshake, more later

# close

- `#include <unistd.h>`
- `int close(int fildes);`
- Initiate a shutdown of connection
  - In-flight data typically finishes
- Free up resources

# struct sockaddr\_in

- struct sockaddr\_in {  
    \_\_uint8\_t        sin\_len;  
    sa\_family\_t      sin\_family;  
    in\_port\_t        sin\_port;  
    struct  in\_addr  sin\_addr;  
    char             sin\_zero[8];  
};
- struct in\_addr {  
    in\_addr\_t        s\_addr;  
};

# Typical usage

```
struct sockaddr_in saddr;  
  
/* Zero out the memory */  
bzero(&saddr, sizeof(saddr));  
  
saddr.sin_family = PF_INET;  
saddr.sin_port = htons(1234);  
saddr.sin_addr.s_addr =  
    htonl(INADDR_ANY);
```

# send/recv data

- `#include <sys/types.h>`  
`#include <sys/uio.h>`  
`#include <unistd.h>`
- `ssize_t read(int fd, void *buf, size_t sz);`
- Read `sz` bytes from `fd` into `buf`
- Similar prototype for `write`
- Pretty much the same for `send/recv`, includes additional `flags` parameter

# Other Useful Functions

- `inet_aton()` and `inet_pton()`
  - `"127.0.0.1"` => sockaddr format
- `gethostbyname()` and `gethostbyaddr()`
  - `"www.google.com"` => hostent
  - `"172.217.10.4"` => `"www.google.com"`
- `getaddrinfo()` and `getnameinfo()`
  - IPv6-friendly versions
- Use the book / man pages / google for more information

# daytimetcpci.c

```
#include "unp.h"

int
main(int argc, char **argv)
{
    int                sockfd, n;
    char               recvline[MAXLINE + 1];
    struct sockaddr_in servaddr;

    if (argc != 2)
        err_quit("usage: a.out <IPaddress>");

    if ( (sockfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        err_sys("socket error");

    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family = AF_INET;
    servaddr.sin_port   = htons(13);    /* daytime server */
    if (inet_pton(AF_INET, argv[1], &servaddr.sin_addr) <= 0)
        err_quit("inet_pton error for %s", argv[1]);

    if (connect(sockfd, (SA *) &servaddr, sizeof(servaddr)) < 0)
        err_sys("connect error");

    while ( (n = read(sockfd, recvline, MAXLINE)) > 0) {
        recvline[n] = 0;    /* null terminate */
        if (fputs(recvline, stdout) == EOF)
            err_sys("fputs error");
    }
    if (n < 0)
        err_sys("read error");

    exit(0);
}
```



# daytimetcpsrv.c

```
#include "unp.h"
#include <time.h>

int
main(int argc, char **argv)
{
    int                listenfd, connfd;
    struct sockaddr_in servaddr;
    char               buff[MAXLINE];
    time_t             ticks;

    listenfd = Socket(AF_INET, SOCK_STREAM, 0);

    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family = AF_INET;
    servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
    servaddr.sin_port = htons(13); /* daytime server */

    Bind(listenfd, (SA *) &servaddr, sizeof(servaddr));

    Listen(listenfd, LISTENQ);

    for ( ; ; ) {
        connfd = Accept(listenfd, (SA *) NULL, NULL);

        ticks = time(NULL);
        snprintf(buff, sizeof(buff), "%.24s\r\n", ctime(&ticks));
        Write(connfd, buff, strlen(buff));

        Close(connfd);
    }
}
```

# Client / Server Model

---

Network Programming

# Clients

- Fairly simple
  - Typically connects to centralized server
- Only concerned with single connection (yours)
- Ex: a web browser
  - One of thousands of users per website
  - Should not impact you at all

# daytimetcpcli.c

```
#include "unp.h"

int
main(int argc, char **argv)
{
    int                sockfd, n;
    char               recvline[MAXLINE + 1];
    struct sockaddr_in servaddr;

    if (argc != 2)
        err_quit("usage: a.out <IPaddress>");

    if ( (sockfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        err_sys("socket error");

    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family = AF_INET;
    servaddr.sin_port   = htons(13);    /* daytime server */
    if (inet_pton(AF_INET, argv[1], &servaddr.sin_addr) <= 0)
        err_quit("inet_pton error for %s", argv[1]);

    if (connect(sockfd, (SA *) &servaddr, sizeof(servaddr)) < 0)
        err_sys("connect error");

    while ( (n = read(sockfd, recvline, MAXLINE)) > 0) {
        recvline[n] = 0;    /* null terminate */
        if (fputs(recvline, stdout) == EOF)
            err_sys("fputs error");
    }
    if (n < 0)
        err_sys("read error");

    exit(0);
}
```

# Servers

- Typically far more complicated than clients
- Have to juggle resources, incoming connections, etc.
  - But often have better CPUs, faster disks, more RAM
- Maintain high uptime (aim for 5 9's)

# Daemons



- Background processes
  - Often initialized during boot
- Network services often under *inetd*
  - Replaces telnet, ftp, www, etc. each having to run their own server
  - Could replace dozens of services + redundant startup code
  - Each would take slot in process table (remember OS?)

# daemon\_init.c

```
#include "unp.h"
#include <syslog.h>

#define MAXFD 64

extern int daemon_proc; /* defined in error.c */

int
daemon_init(const char *pname, int facility)
{
    int i;
    pid_t pid;

    if ( (pid = Fork()) < 0)
        return (-1);
    else if (pid)
        _exit(0); /* parent terminates */

    /* child 1 continues... */

    if (setsid() < 0) /* become session leader */
        return (-1);

    Signal(SIGHUP, SIG_IGN);
    if ( (pid = Fork()) < 0)
        return (-1);
    else if (pid)
        _exit(0); /* child 1 terminates */
}
```

# daemon\_init.c pt. 2

```
/* child 2 continues... */

daemon_proc = 1;           /* for err_XXX() functions */

chdir("/");                /* change working directory */

/* close off file descriptors */
for (i = 0; i < MAXFD; i++)
    close(i);

/* redirect stdin, stdout, and stderr to /dev/null */
open("/dev/null", O_RDONLY);
open("/dev/null", O_RDWR);
open("/dev/null", O_RDWR);

openlog(pname, LOG_PID, facility);

return (0);                /* success */
}
```



# inetd/daytimetcpsrv2.c

```
#include      "unp.h"
#include      <time.h>

int
main(int argc, char **argv)
{
    int listenfd, connfd;
    socklen_t addrlen, len;
    struct sockaddr *cliaddr;
    char buff[MAXLINE];
    time_t ticks;

    if (argc < 2 || argc > 3)
        err_quit("usage: daytimetcpsrv2 [ <host> ] <service or port>");

    daemon_init(argv[0], 0);

    if (argc == 2)
        listenfd = Tcp_listen(NULL, argv[1], &addrlen);
    else
        listenfd = Tcp_listen(argv[1], argv[2], &addrlen);

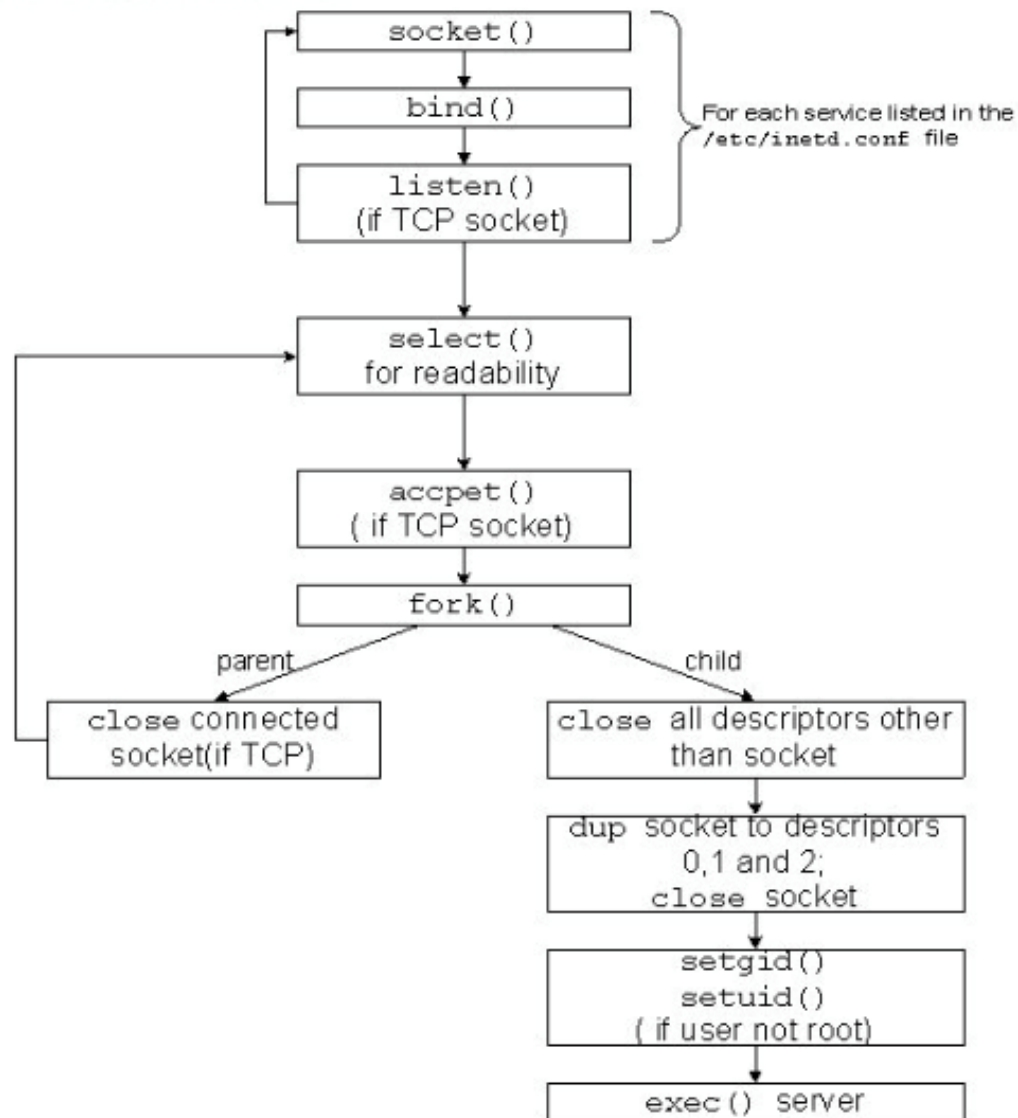
    cliaddr = Malloc(addrlen);

    for ( ; ; ) {
        len = addrlen;
        connfd = Accept(listenfd, cliaddr, &len);
        err_msg("connection from %s", Sock_ntop(cliaddr, len));

        ticks = time(NULL);
        snprintf(buff, sizeof(buff), "%.24s\r\n", ctime(&ticks));
        Write(connfd, buff, strlen(buff));

        Close(connfd);
    }
}
```

# inetd flow diagram



# Server Designs

- Iterative: completely finish work before handling next request
- Concurrent (fork): one process / client
- Concurrent (thread): one thread / client
- Server using **select()**
- Pre-forking: create a pool of processes ahead of time
- Pre-threading: create a pool of threads ahead of time

# Threads



**Ned Batchelder**

@nedbat

 Follow



Some people, when confronted with a problem, think, "I know, I'll use threads," and then two they hav erpoblesms.

RETWEETS

**1,701**

LIKES

**388**



8:47 AM - 23 Apr 2012



32



1.7K



388

# daytimetcpsrv.c

```
#include "unp.h"
#include <time.h>

int
main(int argc, char **argv)
{
    int                listenfd, connfd;
    struct sockaddr_in servaddr;
    char               buff[MAXLINE];
    time_t             ticks;

    listenfd = Socket(AF_INET, SOCK_STREAM, 0);

    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family = AF_INET;
    servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
    servaddr.sin_port = htons(13); /* daytime server */

    Bind(listenfd, (SA *) &servaddr, sizeof(servaddr));

    Listen(listenfd, LISTENQ);

    for ( ; ; ) {
        connfd = Accept(listenfd, (SA *) NULL, NULL);

        ticks = time(NULL);
        snprintf(buff, sizeof(buff), "%.24s\r\n", ctime(&ticks));
        Write(connfd, buff, strlen(buff));

        Close(connfd);
    }
}
```

# tcpserv01.c

```
#include "unp.h"

int
main(int argc, char **argv)
{
    int                listenfd, connfd;
    pid_t              childpid;
    socklen_t          clilen;
    struct sockaddr_in cliaddr, servaddr;

    listenfd = Socket(AF_INET, SOCK_STREAM, 0);

    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family      = AF_INET;
    servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
    servaddr.sin_port        = htons(SERV_PORT);

    Bind(listenfd, (SA *) &servaddr, sizeof(servaddr));

    Listen(listenfd, LISTENQ);

    for ( ; ; ) {
        clilen = sizeof(cliaddr);
        connfd = Accept(listenfd, (SA *) &cliaddr, &clilen);

        if ( (childpid = Fork()) == 0 ) { /* child process */
            Close(listenfd); /* close listening socket */
            str_echo(connfd); /* process the request */
            exit(0);
        }
        Close(connfd); /* parent closes connected socket */
    }
}
```

# select() system call

- Enables handling reading, writing, exceptional cases for file descriptors
- Informs kernel of interest in those specific file descriptors
- ```
int select(int nfd, fd_set *readfds, fd_set *writefds, fd_set *errorfds, struct timeval *timeout);
```

# Clearing / Toggling FDS

- `void FD_SET(int fd, fd_set *s)`
- `void FD_CLR(int fd, fd_set *s)`
- `int FD_ISSET(int fd, fd_set *s)`
- `void FD_ZERO(fd_set *s)`



# Summary

- If server load not heavy, concurrent server with **fork()**
- Pool of servers can “pre-pay” the cost associated with clients connecting
- **select()** can do a lot of heavy lifting for you (more on this next time)

# Lab

- Lab released today at noon, Friday = lab day
- fork() lab
  - Read entire handout
  - Extra Credit (1 lab point) on second page
  - Recommend solving the lab w/o Extra Credit first
- Ask questions when you have them
- Once you're done with the lab make sure to get checked off by a mentor/TA/me
  - Wednesday office hours **this week** end at 11:30am