

TCP State Machine

Network Programming

Overview


- TCP (Transmission Control Protocol) provides reliability
 - In-order
 - Connection-oriented
- Uses a "state machine"
 - transitions from one state to another based on input / actions

IP

- Connectionless delivery
 - Each datagram handled individually
- Unreliable
 - No guarantee
- Fragmentation + reassembly
 - Hardware MTU
- Routing
- Error detection

TCP Segment Format


<https://tools.ietf.org/html/rfc793#section-3.1>



TCP Header Format

Note that one tick mark represents one bit position.

Three Way Handshake



Three Way Handshake

- Initiator sends first segment
 - SYN, seq=x
- Listener responds
 - SYN, seq=y, ACK=x+1
- Initiator ACKs the ACK
- Necessary and sufficient

TCP Builds on IP

- Unreliable packet delivery
- Messages can be lost, delayed, duplicated, or arrive out of order
- TCP requires a timeout mechanism
 - If we don't receive an ACK for segments we have sent, resend

3WHS: What We Get

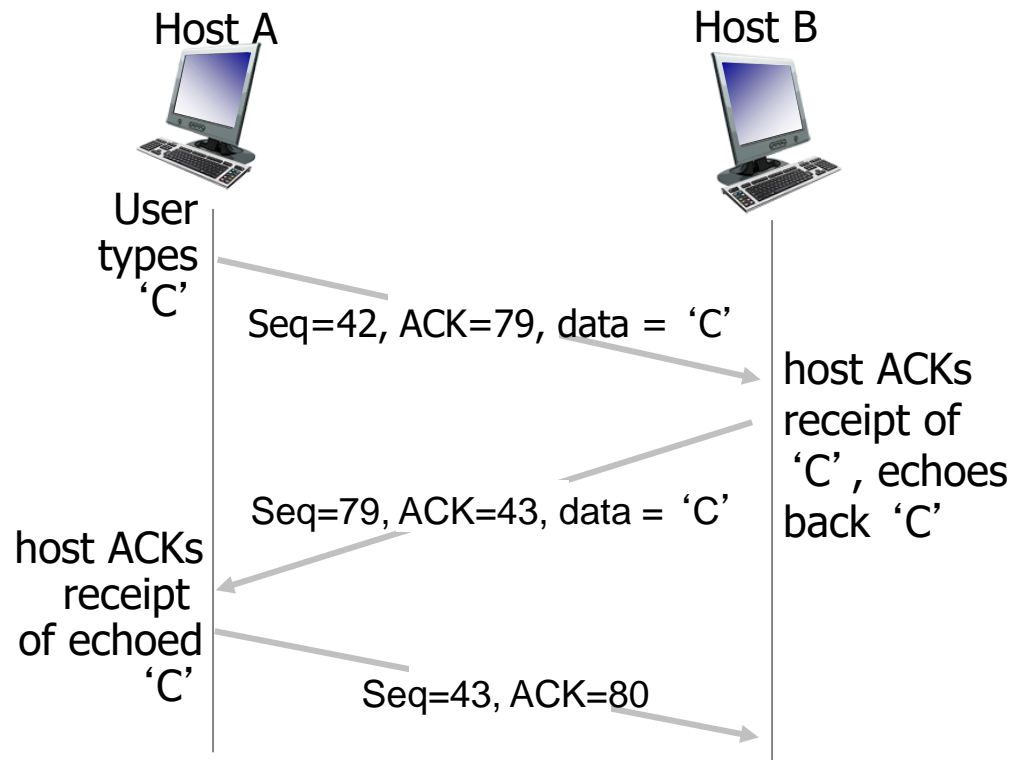
- Guarantees both sides are ready
 - Also, they know they are ready
- Both have agreed to their respective Initial Sequence Numbers
- RFC 761, Section 3.3 covers ISNs
 - <https://tools.ietf.org/html/rfc761#section-3.3>
 - “It should be noted that this strategy does not protect against spoofing or other replay type duplicate message problems.”

TCP "Normal" Behavior

- Send the data
- Wait for an acknowledgement (ACK)
- More data? Repeat.
- Simple, right?

Simple Telnet Example

- Source on last slide



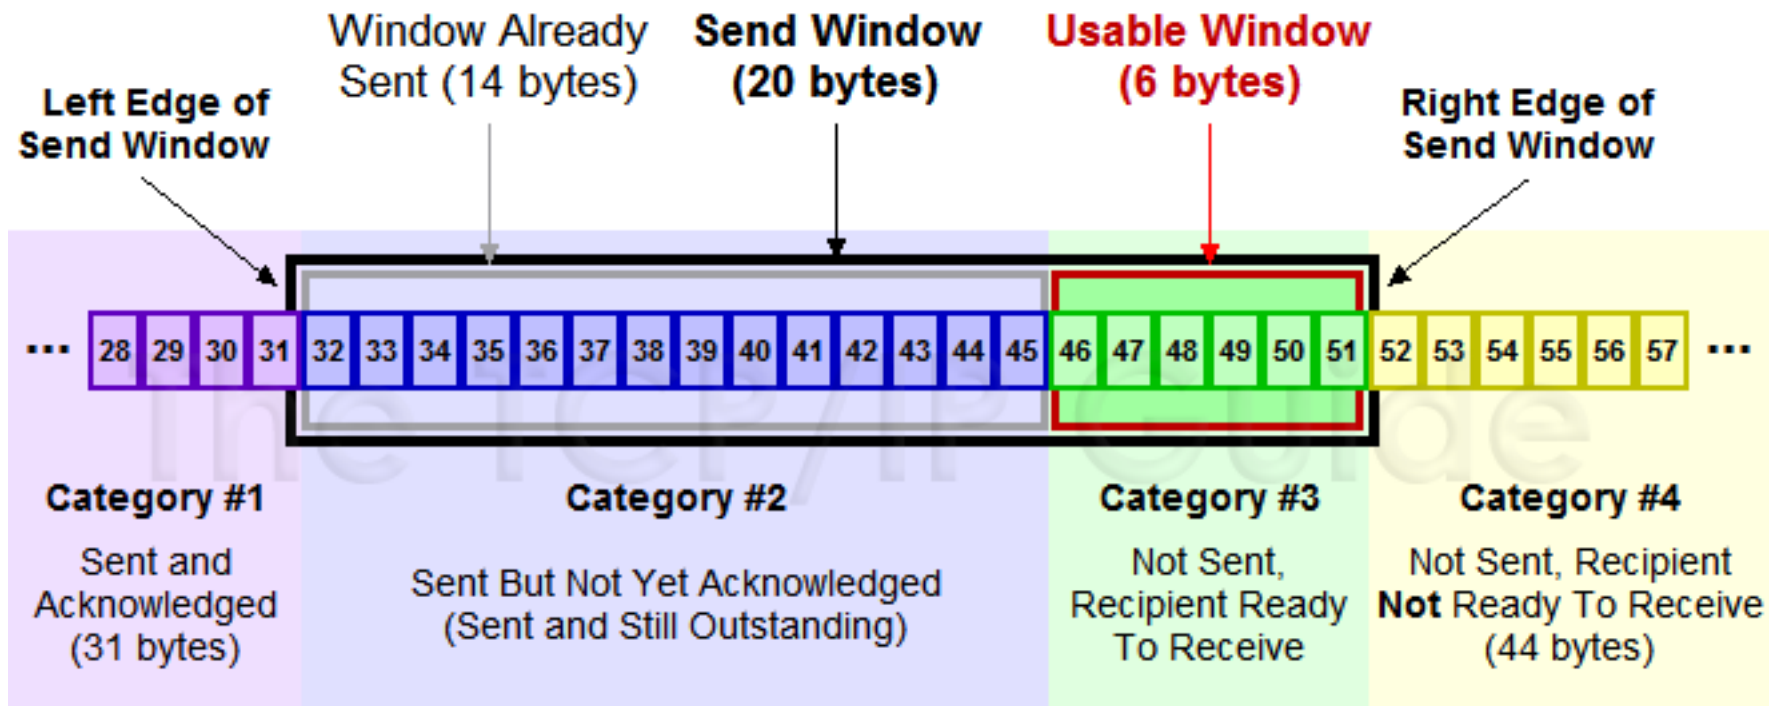
simple telnet scenario

High latency networks?

- We would spend a lot of time waiting for ACKs
- The network would go largely underutilized
- Any way around this problem?

Sliding (dynamic) Windows

- <http://packetlife.net/blog/2010/jun/7/understanding-tcp-sequence-acknowledgment-numbers/>



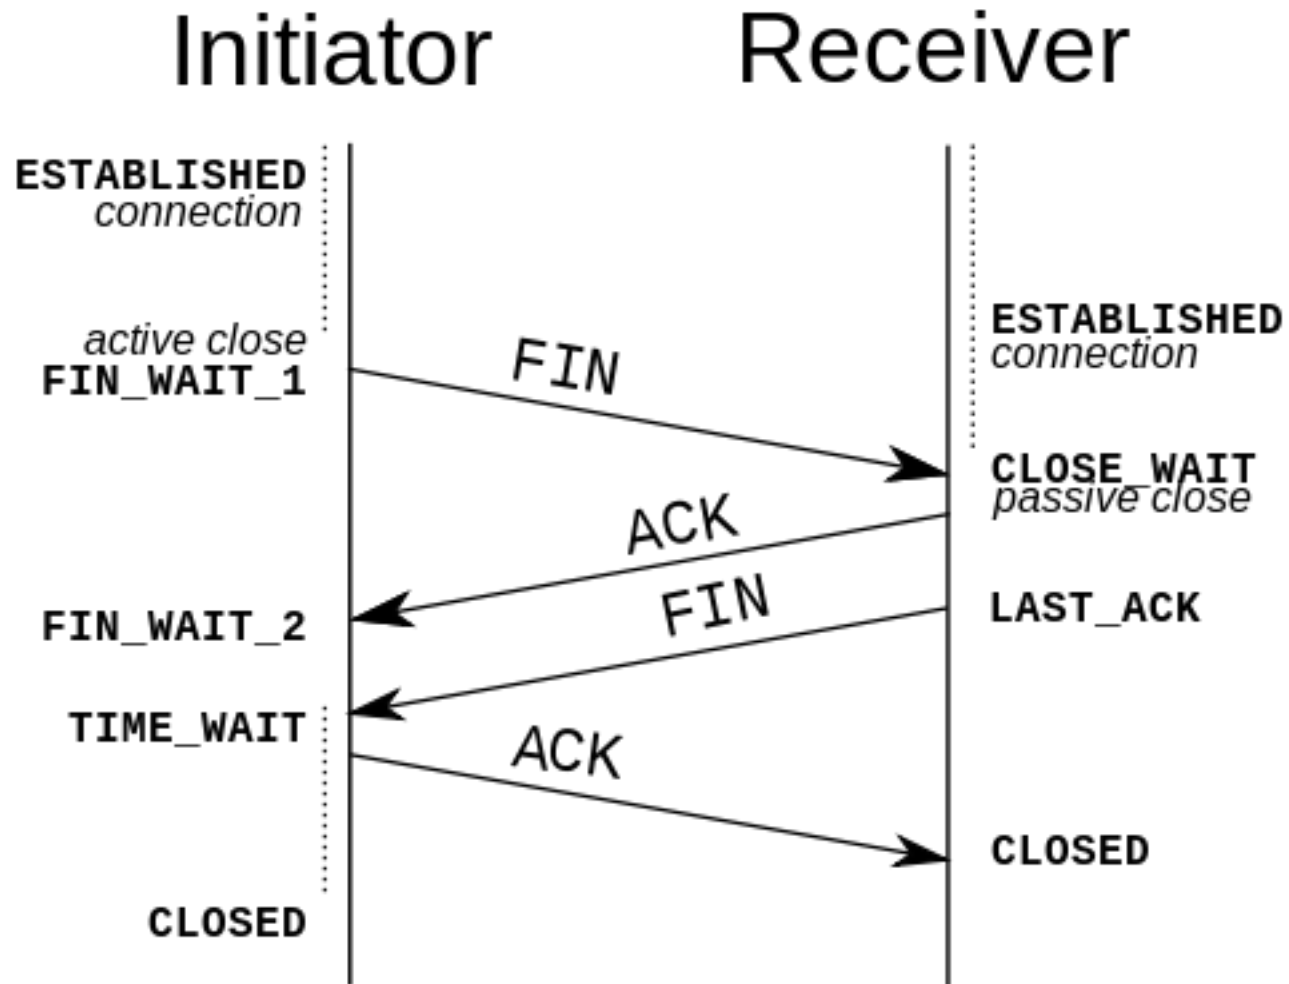
TCP Close Connections

- Use a modified 3WHS
- Recall TCP is full-duplex, there are two independent streams
- When A decides it's done, close connection (one direction):
 - wait for final ACK, then send FIN
 - B ACKs the FIN

TCP Close Connections pt 2

- B decides to close its connection
 - Send FIN, ACK $x+1$
- A receives FIN + ACK segment
 - Send ACK $y+1$
- Connection closed in both directions

Four Way Handshake



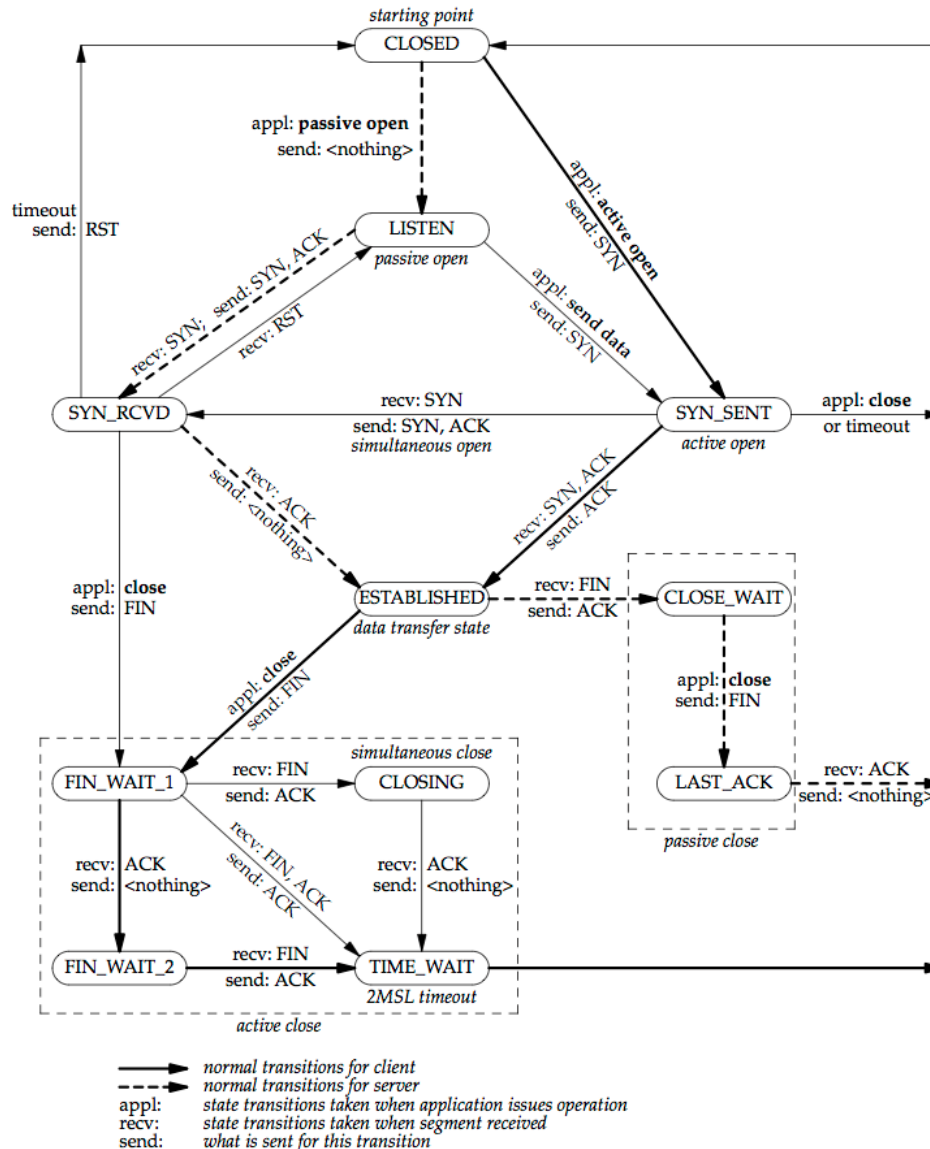
4WHS: Details

- Initial FIN is important
 - TCP ACKs and informs application
 - Receiver FIN _NOT_ sent immediately
- Receiver app closes connection
 - Sends 2nd FIN, ACK
- Initiator receives FIN+ACK, responds with ACK

TCP Connection Reset

- Normally use `close()` to close connection
- Sometimes things go wrong; we must break the connection
- Send a TCP segment with `RST=1`
- Other host aborts connection

TCP State Machine



Source:

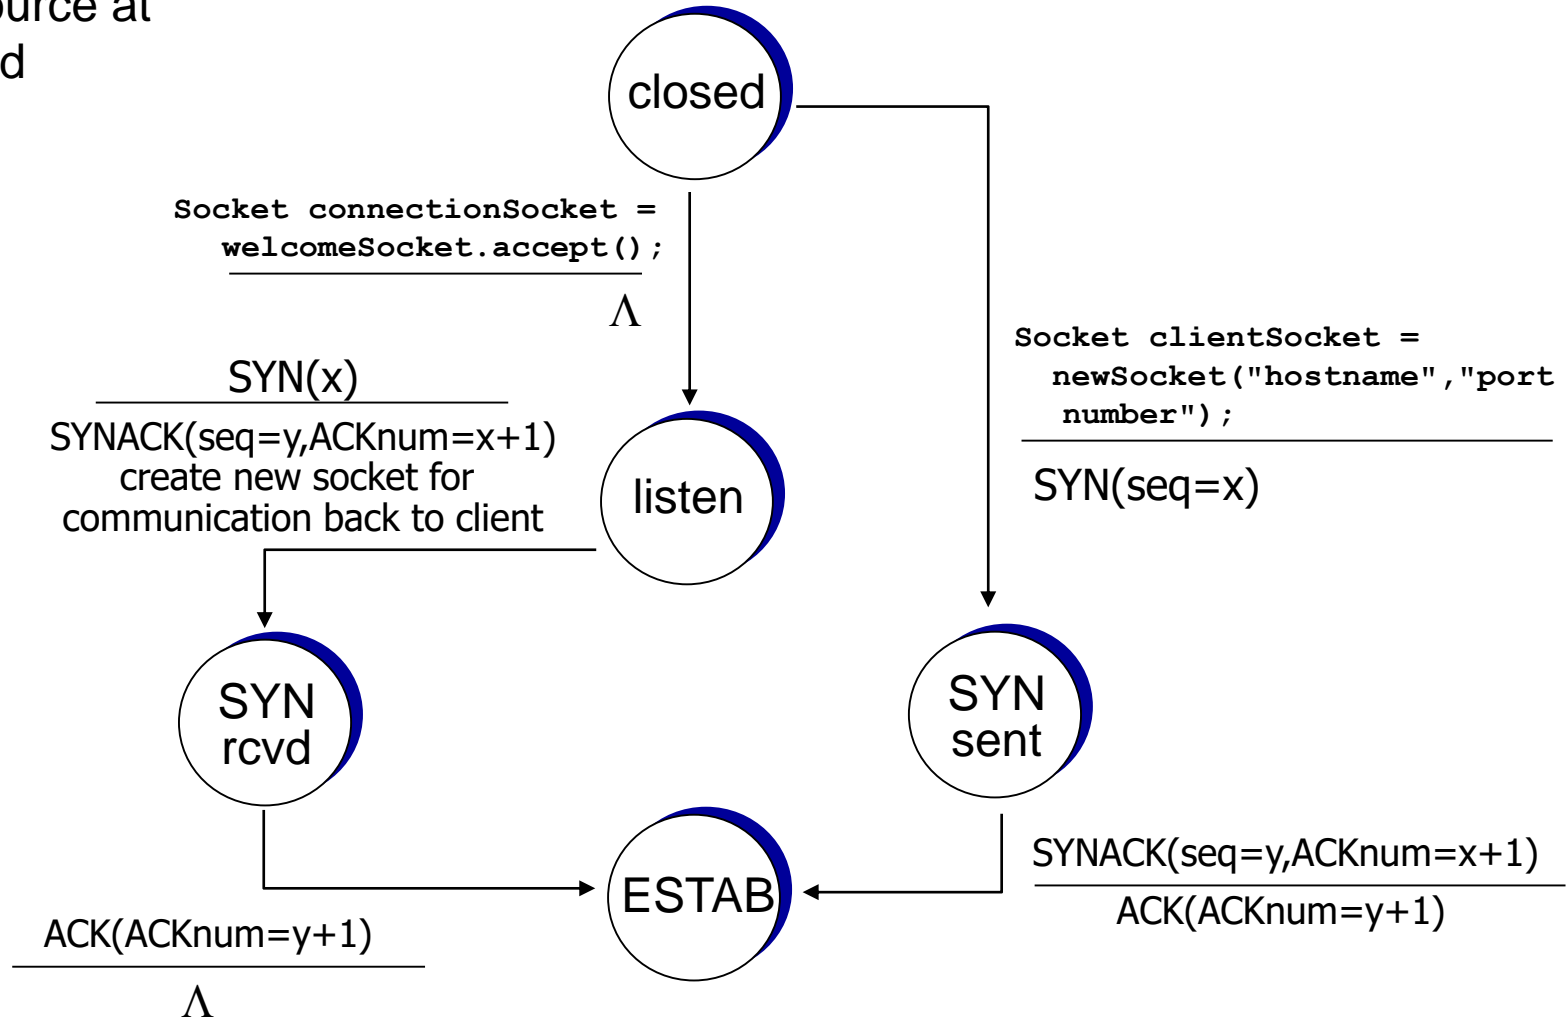
https://www.ibm.com/support/knowledgecenter/en/SSLTBW_2.1.0/com.ibm.zos.v2r1.halu101/constatus.htm

TCP State Machine

- Begins in CLOSED state
- Passive open (wait for connection)
- Active open (initiate connection)
 - When moving from CLOSED to SYN SENT, emit SYN segment
 - Other host sends SYN/ACK, send ACK and move to ESTABLISHED

TCP 3-way handshake: FSM

Source at
end



TCP: closing a connection

client state

ESTAB

`clientSocket.close()`

FIN_WAIT_1

can no longer
send but can
receive data

FIN_WAIT_2

wait for server
close

TIMED_WAIT

timed wait
for $2 * \text{max}$
segment lifetime

CLOSED



FINbit=1, seq=x

ACKbit=1; ACKnum=x+1

FINbit=1, seq=y

ACKbit=1; ACKnum=y+1

can still
send data

can no longer
send data

server state

ESTAB

CLOSE_WAIT

LAST_ACK

CLOSED

Source at
end

TIME_WAIT State

- Handles unreliable delivery
- TCP defines MSL for transit lifetime
- Don't let old streams interfere with current streams
 - Timers allow TCP to distinguish between old and new streams
 - Prevents old duplicates from being mistaken for segments in current stream

HW2 Clarification

- Usernames, guess words, and secret words are **not** case-sensitive

Chapter 3

Transport Layer


A note on the use of these Powerpoint slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016
J.F Kurose and K.W. Ross, All Rights Reserved



Computer Networking: A Top Down Approach

7th edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016