

# TCP Congestion Control

---

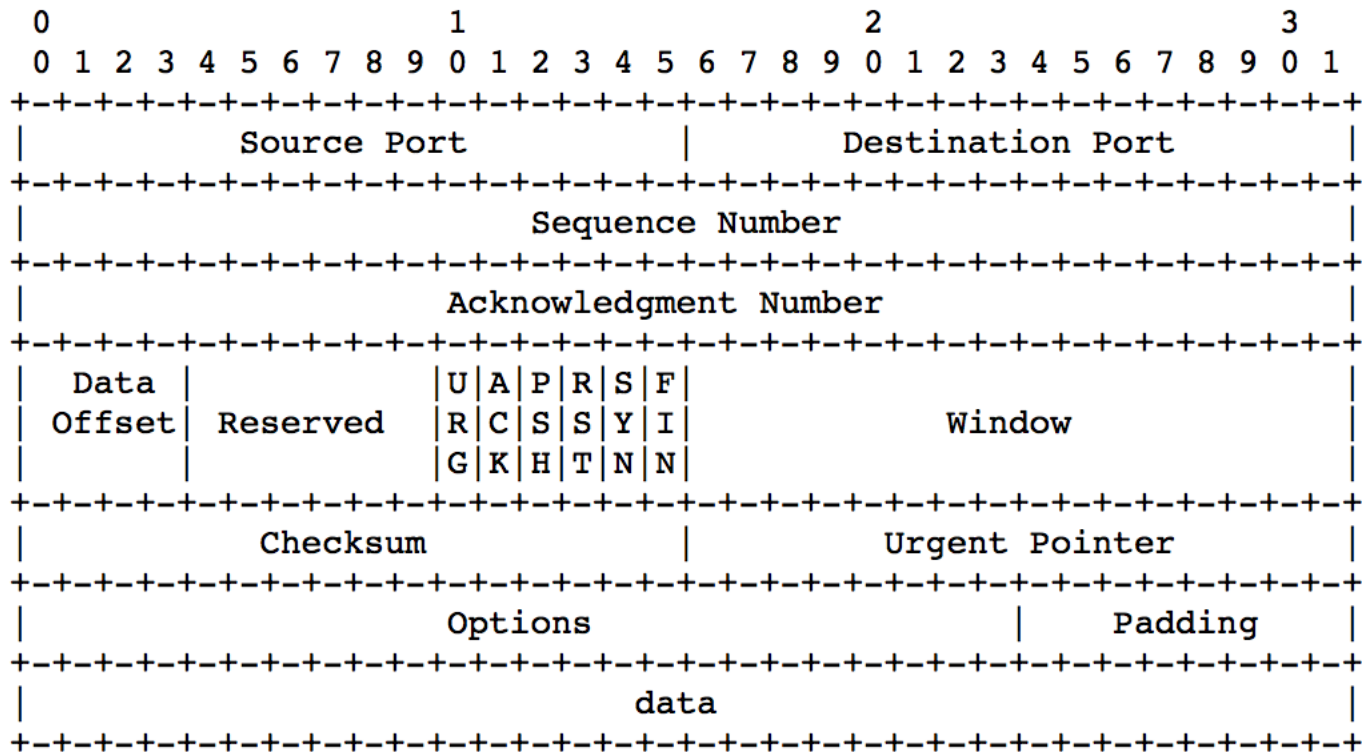
Network Programming

# Overview

- TCP (Transmission Control Protocol) provides reliability
  - In-order
  - Connection-oriented
- Also provides two important items:
  - Flow Control
  - Congestion Control

# TCP Segment Format

<https://tools.ietf.org/html/rfc793#section-3.1>



TCP Header Format

Note that one tick mark represents one bit position.

# Flow Control

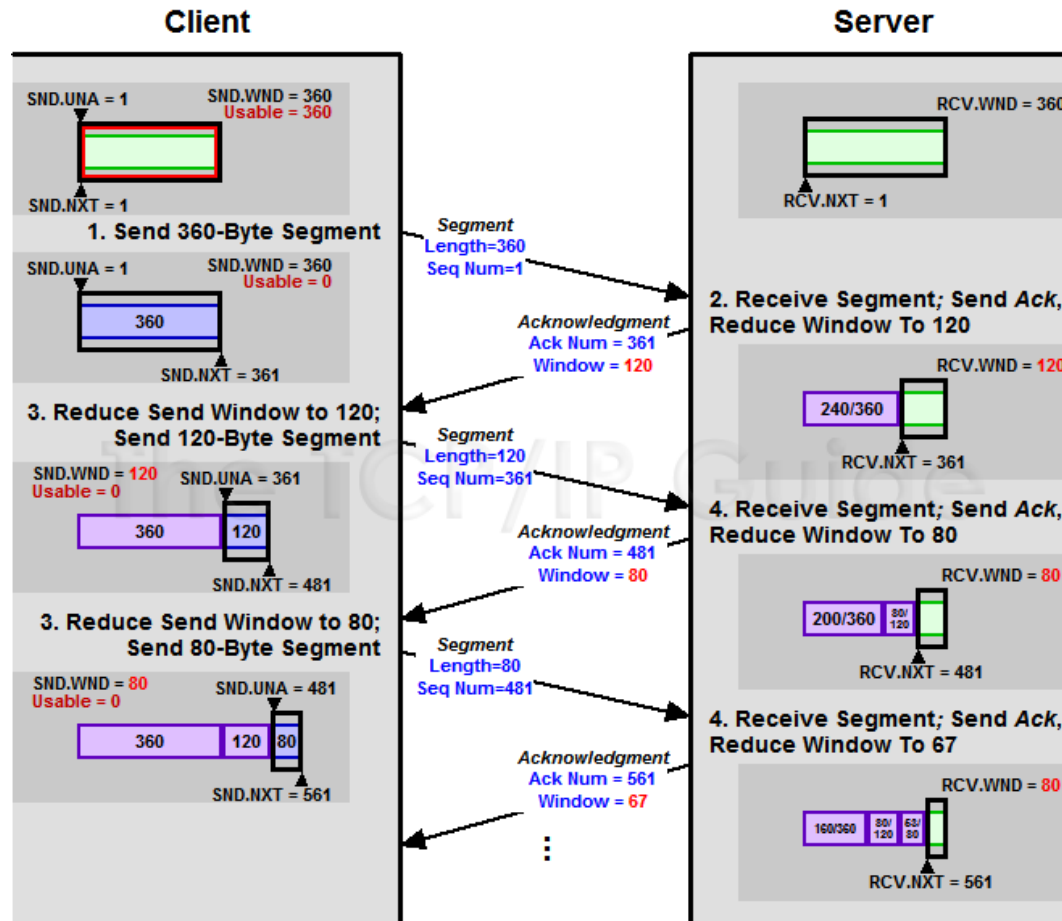
- Endpoints today are fast!
- Can easily saturate most connections
  - In turn, can easily saturate receiving hosts
- Requires a mechanism to tell sending host to slow down

# Sliding Window

- Allows receiver to send feedback to sender
- Communicates *receive window* for bytes it's willing to accept
  - If 0, can't send at all!
- Sender cannot send more than advertised window space

# Silly Window Syndrome

- Slow consumer / producer problem



# Nagle's Algorithm

- With small interactive sessions, small (1 byte?) data transmissions
- With a 40 byte TCP header, huge overheads!
- Combine these small messages
  - Send all at once, achieve greater efficiency

# Congestion Control

- Congestion leads to packet loss
  - Increases retransmissions
- TCP generally follows pattern of Additive Increase, Multiplicative Decrease (AIMD)
  - Basically, accelerate slowly but brake hard
- Force sender to slow down



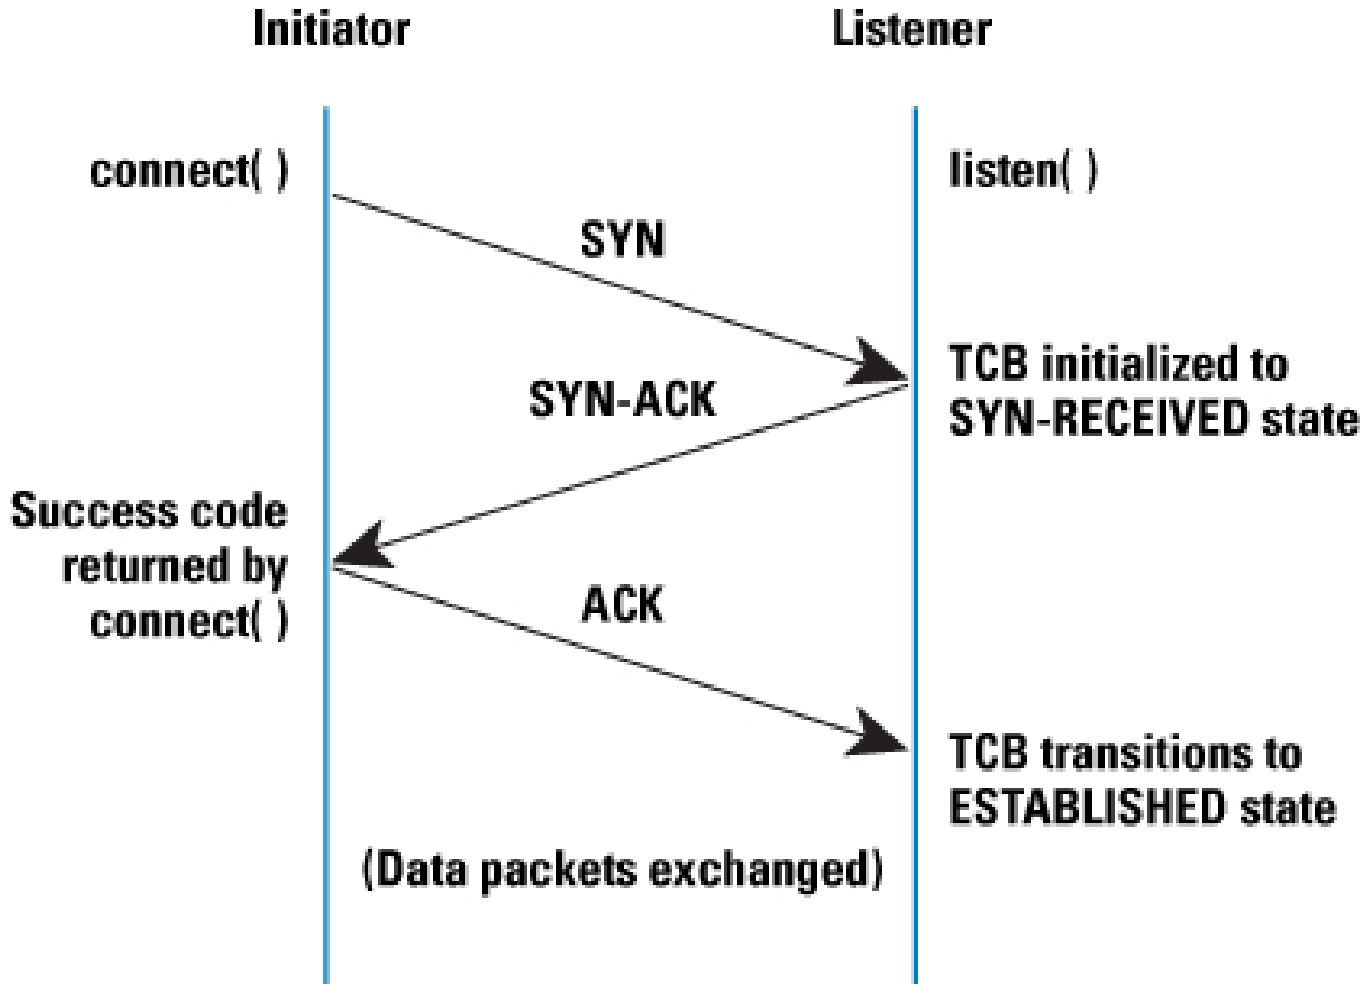
# Congestion Collapse

- If left unchecked, could destroy useful bandwidth of network
- In 1986, connection between LBL and UC Berkeley (400 yards) dropped to 40 bps due to this problem
  - Factor of ~1000 drop

# CC Countermeasures

- <https://tools.ietf.org/html/rfc5681>
- Slow Start
- Congestion Avoidance
- Fast Retransmit
- Fast Recovery

# Three Way Handshake



# Slow Start

- Initialize congestion window to one segment (MSS)
  - Established during 3WHS
- On each ACK increase cwnd
  - $\text{cwnd} \leftarrow \text{cwnd} + \text{MSS}$
- Effect? Exponential growth of cwnd during each RTT
  - Named for initial size, not slow!

# Congestion Avoidance

- Slow Start initially used
  - But packet gets dropped
  - Or slow start threshold (ssthresh) exceeded
- Network congestion is happening
  - Reduce to linear increase
- On each ACK increase cwnd
  - $cwnd \leftarrow cwnd + MSS^2/cwnd$

# Fast Retransmit

- Duplicate ACK received. Why?
  - Lost? Delayed?
- If three or more duplicate ACKs are received, immediately retransmit last segment
  - Bypass timers

# Fast Recovery

- Duplicate ACKs can only be generated when a segment is received
  - Network may not be as bad as imagined
- Enter Congestion Avoidance mode instead of resetting window size by jumping into Slow Start

# RED / ECN

- Random Early Discard takes place at burdened routers
  - As burden goes up, so does likelihood of RED happening
  - Triggers multiplicative decrease
- Explicit Congestion Notification
  - RED may be too extreme
  - Routers toggle bit in header, receiver echos it, sender slows transmission



# TCP Congestion Avoidance Flavors

---

Network Programming

# Slightly Deeper Look...

- We just discussed TCP congestion control
- One of the big things TCP provides is Flow Control
- We said “AIMD” and left it at roughly that.
  - Is it that simple?

# Evolving Options

- The basic principle is the same:
  - Get to about the right amount of traffic
  - React to congestion appropriately
  - Increase traffic
- The overall goal is to minimize
  - Underutilization of bandwidth
  - Flooding of hosts / internetwork (i.e. dropped packets)

# Wikipedia To the Rescue

- Great table

# TCP Tahoe

- If 3 duplicate ACKS are received
  - Things are really bad!
- Fast Retransmit
- $ssthresh = cwnd/2$
- $cwnd = 1 \text{ MSS}$
- Slow Start

# TCP Reno

- If 3 duplicate ACKS are received
  - Things are really bad! (Again)
- Fast Retransmit
- $\text{tcp.cwnd} *= 0.5$
- $\text{ssthresh} = \text{cwnd}$
- Fast Recovery

Next 2 slides are courtesy of...

# Chapter 3

## Transport Layer


### A note on the use of these Powerpoint slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016  
J.F Kurose and K.W. Ross, All Rights Reserved



## *Computer Networking: A Top Down Approach*

7<sup>th</sup> edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016


# TCP: switching from slow start to CA

**Q:** when should the exponential increase switch to linear?

**A:** when **cwnd** gets to 1/2 of its value before timeout.

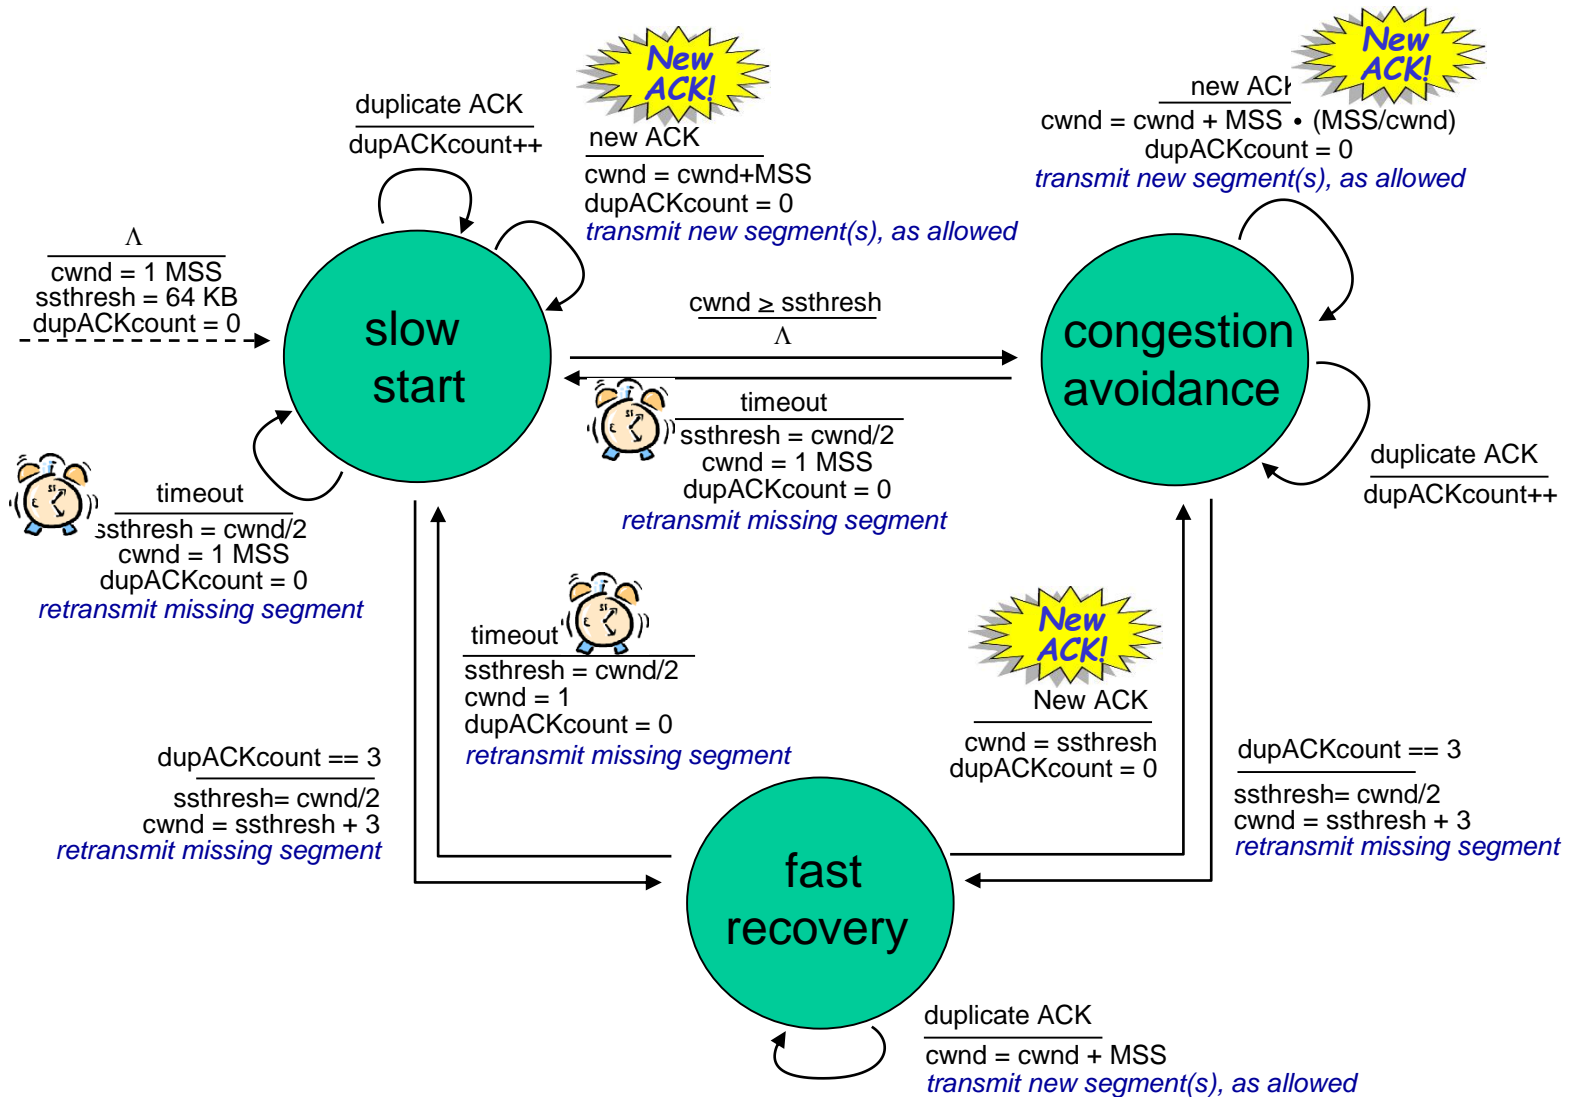
## Implementation:

- variable **ssthresh**
- on loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event





# Summary: TCP Congestion Control



# Some More...

- TCP Vegas
  - Use RTT changes to detect congestion
  - Not very widespread
- TCP Westwood
  - Like Reno but use bandwidth estimation instead of just cutting the window
  - EWMA, other “low pass” solutions?
- TCP New Reno
  - During Fast Recovery
  - Any duplicate ACK yields a new packet

# TCP BIC/CUBIC

- Used in Linux Kernel
  - BIC 2.6.8 – 2.6.18
  - CUBIC 2.6.19+
- BIC is built for LFNs
- CUBIC uses a cubic “smoothing”
  - Like BIC but doesn’t grow as fast