

HW3 Hint

- In this assignment we are not passing in an IP address, rather we are passing in a hostname
- As a result, you will need to apply what you learned in Lab 6 in order to get an address that you can use to `connect()` to the server
- Keep in mind that `connect()` expects an address in network order so you should not be calling `inet_ntoa()` or `inet_ntop()` when populating a `struct sockaddr`.

gRPC

Network Programming

RPC Definition

- Remote Procedure Calls (RPCs)
 - Should appear like a local function call
 - Code executed by call may be on another host
 - Remote code should return when call is complete, possibly pass back data
- RPCs use message passing as the underlying mechanism
- The goal is to make calls instead of requesting resources like in Representational State Transfer (REST)

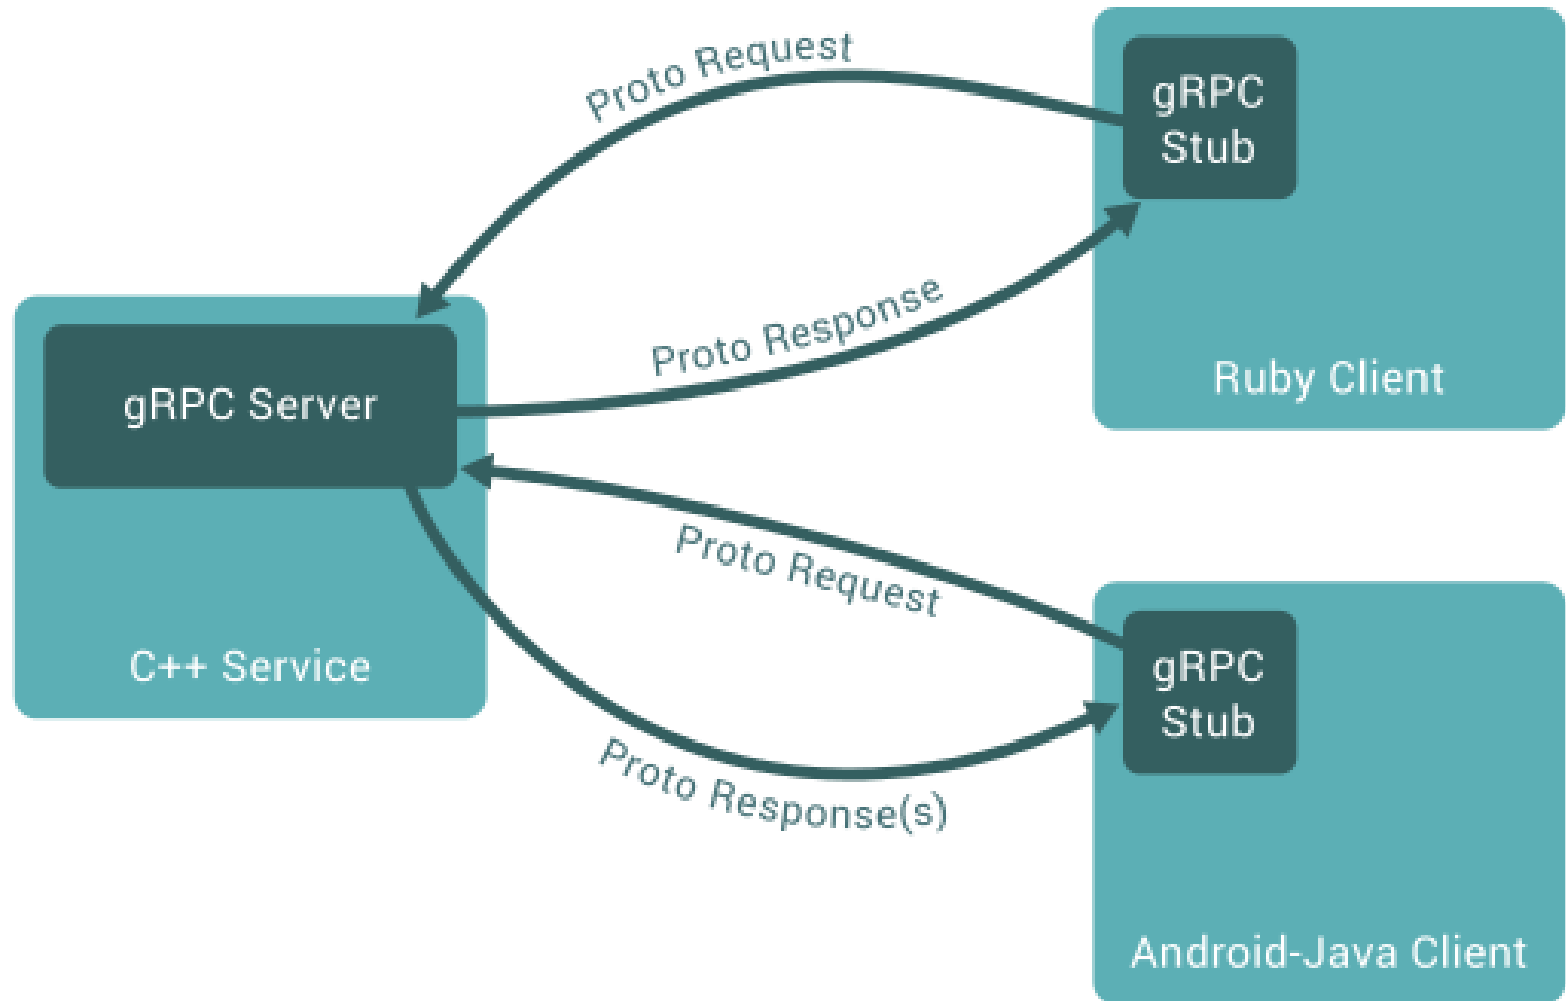
Why Use RPCs?

- RPC protocols hide network details
- Simplifies code, lets us write code as though we were using local functions
- Offload responsibilities to specialized hosts
- Maybe...
 - Only certain hosts have a database engine
 - We don't want to replicate data
 - Some hosts are better at doing fast arithmetic
 - Etc.

Brief History

- 1980s – Practical network implementations
- 1991 – Common Object Request Broker Architecture (CORBA)
- 1998 - XML-RPC
- 2000 - Simple Object Access Protocol (SOAP)
- 1999 - XMLHttpRequest developed
- 2002-2007 – Various browsers officially add XMLHttpRequest
- 2015 – gRPC released by Google

gRPC Flowchart



gRPC proto3 (protocol buffers)

- Note that *pb2* which will show up in some code is for Python Protocol Buffers API v2. This works with proto2 or proto3 (protocol buffer description languages)
- Describes *service*, *rpc*, and *message* types
 - Full documentation online but .proto file is human readable
 - Real data is compressed/serialized and not human readable. **Can** be much more efficient than XML
- *messages* typically correspond to OOP objects

gRPC Channels

- Communication between a server and a client happens over a *channel*
- `with`
`grpc.insecure_channel('localhost:50051')`
`as channel:`
 - The *with* is a context manager in Python – we'll close the channel when we go out of scope
 - We will need to give our *stub* a channel to communicate over (more on this soon)
 - This example uses an unauthenticated (insecure) channel, easy and works for us

gRPC Authentication

- Much like HTTP vs HTTPS, insecure_channel isn't always sufficient
- SSL/TLS support
 - Need a certificate
- Google tokens via OAuth2
- Can have security over a channel
- Can also do per-call security
- Examples:
<https://grpc.io/docs/guides/auth.html#python>

gRPC stubs

- Clients use stubs to actually invoke the RPC calls
- A stub looks like the RPC prototype but only sends one or more messages, and then receives one or more messages back.
- As far as the client can tell, we made a call and got the expected output back
 - Assuming no errors

gRPC Errors

- Calls can raise errors, in Python the client may print runtime errors that are from the server side
- General errors such as
GRPC_STATUS_CANCELLED,
GRPC_STATUS_UNIMPLEMENTED
- Network errors such as
GRPC_DEADLINE_EXCEEDED
- Protocol errors such as
GRPC_STATUS_UNAUTHENTICATED

gRPC Timeout/Deadline

- Depending on language, you may have a timeout (number of seconds) or a deadline (absolute timestamp)
- By default there is no timeout/deadline and calls can run forever
 - We'll use this default in lab, but in general you should always have timeouts to be safe
- In Python:
 - `helloworld_pb2.HelloReply(message='Hello, %s!' % request.name, timeout=2)`

gRPC Method Types

- Unary

`rpc y(InMessage) returns (OutMessage) {}`

- Client Streaming

`rpc y(stream InMessage) returns (OutMessage) {}`

- Server Streaming

`rpc y(InMessage) returns (stream OutMessage) {}`

- Bidirectional Streaming

`rpc y(stream InMessage) returns (stream OutMessage) {}`

gRPC Wire Representation

- Uses HTTP2
- Not a casual read: [[docs](#)]
- You can look at a trace in Wireshark
 - Strings tend to be left intact
 - Field names and other fields may not be human readable due to compression / serialization
 - Wireshark won't be able to dissect the application layer data

gRPC Generating Files

- See `examples/python/route_guide/run_codegen.py`
- If we have `x.proto` which describes a service `A` with `rpc y (Z)` where `Z` is a message type containing just a single string `b`
- We will need a stub, for example
`stub = x_pb2_grpc.AStub(channel)`
- Then we can call
`stub.y(x_pb2.Z(b="somevalue"))`
- Creates `Z` according to `x.proto`, uses it in RPC call `y` over channel associated with `stub`

gRPC Servicer Registration

- Server has to implement the RPC calls
 - “compilation” builds a “Servicer” for each `class` defined in the `.proto`, so in our example `AServicer`
 - Always just “Servicer” after the class name
 - Derives from `x_pb2_grpc.AServicer`
 - `class MyClass (x_pb2_grpc.AServicer) :`
- A Servicer needs a server, see examples for `grpc.server()`
- `x_pb2_grpc.add_AServicer_to_server(MyClass() , server)`

Hello World Example - .proto

- examples/protos/helloworld.proto
- Ignore all the option java/objc stuff, we don't need it in Python
- Greeter service has SayHello method
 - Expects a HelloRequest
 - Returns a HelloReply
- **syntax = "proto3";**
 - Important, says we're using v3 of Protocol Buffers

Hello World Example - .proto

- Each of these types is a message
 - Note that using "message" as a variable name is fine (terrible style though)!
 - The assignment shows a unique “tag”, we must enumerate each variable and they cannot share tags
 - You can nest messages in messages
 - Frequently used variables should have tags < 16

Hello World Example - Client

- `examples/python/helloworld/greeter_client.py`
- Have to import `grpc`, `x_pb2`, `x_pb2_grpc` (in this case `x` is `helloworld`, since we used `helloworld.proto`)
- Insecure channel on TCP port 50051 used
 - When the `with` context manager is closed, connection is closed automatically
 - Stub takes channel, follows naming rules
 - Use the stub to call RPC `SayHello`

Hello World Example - Server

- `examples/python/helloworld/greeter_server.py`
- Note that Greeter extends `helloworld_pb2_grpc.GreeterServicer`
 - Didn't need to name our class Greeter
- Server is allotted 10 threads in this example
- `server.start()` does not block, busy wait loop with 1 second sleep, CTRL+C to quit
 - Without some sort of busy wait, after `.start()` we would immediately terminate the server

Synchronous / Asynchronous Calls

- Most of the examples are synchronous
- Sometimes you might want asynchronous calls
- This can be done using Python's native features
- Just use futures (`concurrent.futures` or `asyncio.Future`)

Synchronous / Asynchronous Calls

- Example:

```
call_future =  
stub.SayHello.future(helloworld_pb2.HelloRequest(n  
ame='you'))
```

```
call_future.add_done_callback(process_response)
```

- Source: <https://github.com/grpc/grpc/issues/16329>

Route Guide Example Dataset

- `examples/python/route_guide/route_guide_db.json`
- Contains JSON notation (standard format, easy to parse)
- Array of associative arrays (Features)
- Each Feature associates a "name" with a "location" (Point)
- Each Point is an associative array with an integer "latitude" and "longitude"

Route Guide Example Helpers

- `examples/python/route_guide/route_guide_resources.py`
- Uses the Python json parsing library
- Single function, `read_route_guide_database()`
 - Creates a list which is returned
 - Every entry is a Feature (name (string) + location (Point)), uses the `route_guide_pb2` classes
 - These are based on the message types described in the `.proto`
 - `examples/proto/route_guide.proto`

Route Guide Example .proto

- Single service, RouteGuide
 - GetFeature(Point), Unary
 - ListFeatures(Rectangle), Server Stream
 - RecordRoute(stream Point), Client Stream
 - RouteChat(stream RouteNote), Bidirectional Stream
- Messages
 - Point (int32, int32) - lat. lon.
 - Rectangle(hi, lo) - two corners to make a diagonal
 - Feature(name, Point)
 - RouteNote(location, message) - used in RouteChat
 - RouteSummary(int32 x5) - RecordRoute() results

Route Guide Example Server

- Main loop looks the same as helloworld example
- `get_feature()` returns a `Feature` or `None`
- `get_distance()` calculates distance between two points
- Everything else in `RouteGuideServicer`

Route Guide GetFeature() Server

- `RouteGuideServicer.__init__()` loads the json into `self.db`
- `GetFeature()`
 - `context` contains info like deadlines
 - Input `request` is a `Point` (see `.proto`)
 - Returns `Feature` from our database or `None` if it can't be found

Route Guide GetFeature() Client

- Two calls to `guide_get_one_feature()`, one using a specific Point, one using the origin (0,0)
- `guide_get_one_feature()`
 - Make the gRPC call
 - Sanity check for formatting
 - If there's a name, lookup succeeded, print the name and location
 - Otherwise print that lookup didn't find the location, and print location

Route Guide ListFeatures() Server

- ListFeatures()
 - Input `request` is a Rectangle (see .proto)
 - Client will expect a stream of Features
 - Returning all Features at once could be a lot of data, instead, this is a Server Streaming call
 - Treat the server as a generator by using the `yield` keyword
 - Code runs until first time `yield` is hit, returns `yield` argument
 - Every time after, resume from `yield`, run until `yield/function` end

Route Guide ListFeatures() Client

- `guide_list_features()` creates a `Rectangle` and uses the stub to RPC invoke `ListFeatures()`
- for feature in features:
 - Print each feature
 - Treats the result as an iterable or generator
 - We can start this loop as soon as we get a single Feature back from the server
 - Don't need to wait for all N Features, don't need to store N Features in memory!

Route Guide RecordRoute() Server

- RecordRoute()
 - Input `request_iterator` is a stream of Points (see .proto)
 - Just like the last stream we looked at, there's a for loop
 - Server treats this like a generator, can start as soon as one Point is received, keeps going until stream ends.
 - Do a lot of statistic counting, including distance, how many Points were Features, total length of time the call takes (on localhost this will likely be 0 seconds)
 - Package up all our results into one RouteSummary object, return

Route Guide RecordRoute() Client

- Starts with `guide_record_route()`
- Uses `generate_route()` to get a generator of random Features
- `route_iterator` stores an instantiation of the `generate_route()` generator
- Use the stub to call RPC `RecordRoute()`
 - It's fine that we directly pass a generator
- Print the `RouteSummary` we get back

Route Guide RouteChat() Server

- RouteChat()
 - Input `request_iterator` is a stream of RouteNote (see .proto)
 - Keep a list of previous notes
 - If the latest RouteNote from the stream matches the location of a note in previous notes, return the match from the past as a RouteNote using yield
 - We are returning a stream of RouteNotes as well
 - If we haven't encountered this location, add the new RouteNote to the end of the previous notes list

Route Guide RouteChat() Client

- `guide_route_chat()` gets a `RouteNote` generator `generate_messages()` and passes it directly to `RPC RouteChat`
- Prints every message we receive from the RPC, looping over the return value (because it's also a generator)
- `generate_messages()`
 - Fixed array of `RouteNotes` each made by passing arguments to `make_route_note()`
 - **yield** one message at a time, act as generator

Lab 7

- gRPC lab
- Should have to change relatively little code
- Downloads will take some time, need about 300MB of space between packages and source code
- Refresh yourself on Python3 syntax for printing, functions, objects
 - Remember that globals are frowned upon and to refer to a global inside a function you have to use the **global** keyword

Lab 7

- Both server and client will run on localhost for simplicity, but you could stick the server anywhere and your client should work the same (just with a different remote address in the channel)
- `grpc/examples/python/route/run_codegen.py` will make your life easier. `-I` specifies directory to find the `.proto` in, the `--out` directories being `“.”` (current directory) puts all generated code inside the directory you call `run_codegen` from.