



Peer-to-Peer Networking

Network Programming

Interesting Approaches

- High-level overview
 - Only scratch the surface!
- Too many to cover any in depth
 - High level overview
- I encourage everyone to dig more deeply if anything intrigues you

Motivation

- Routing resilience
 - If a node breaks, should be able to route around the damage
- Communication still feasible
- Formerly, prioritize connectivity
- Now prioritize content from endpoints!

Problems

- What if an endpoint goes down?
 - Gmail? Facebook? Twitter?
- What if a state-level agency decides Facebook is bad?
 - Maybe Bhutan bans depression sites?
- What about natural disasters?
 - I need to find my family/friends

Peer-to-Peer

- All peers are equal
 - Content providers
 - Routing partners
- Pay to play
 - Each host generates workload
 - Each host also contributes resources
- Conceptually, scales well


Examples

- Napster
 - P2P w/ a centralized server
 - Single point of failure!
- Gnutella
 - Creates an overlay network
 - Floods content requests
 - No single point of failure, more resilient than Napster but scalability issues
- Both have limitations!

Distributed Hash Tables

- Properties:
 - Decentralized
 - Fault Tolerant
 - Scalable
- Generally use keyspace partitioning scheme
- Each node maintains a (partial) routing table for the overlay network

Routing in Overlay Network



Gnutella

- Nodes are called servants
- Connect, then send descriptors (Gnutella protocol concept, not C file descriptor)
 - Ping to discover
 - Reply with one or more pongs to describe availability
 - Query to ask for data
 - QueryHit to reply that you have the data

Gnutella Routing

- Pongs can only travel back along the path the Ping came from
- QueryHit similarly can only follow Query's path
- Forward Ping/Query to all direct connections except the one you received it on
- Decrement TTL field and increment Hops before forwarding. Do not forward if $TTL = 0$
- Don't forward if payload descriptor and descriptor ID already seen - prevents replication of messages!

Gnutella Downloads

- Remember that Gnutella is an overlay
- For downloads we do a direct client-to-client connection
 - Don't route through other servants in the overlay network
- This is still P2P, no dedicated server
 - But if we could stop the initial distributor before copies were shared, we could prevent a file from spreading
 - Easier said than done!

BitTorrent

- Break files up into *pieces*
- Upon downloading a piece, able to then serve it to another consumer
 - Popularity increases number of hosts for that piece / content
- Cryptographic hashing used to ensure no changes to piece made

Chapter 2

Application Layer

A note on the use of these Powerpoint slides:


We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016

J.F Kurose and K.W. Ross, All Rights Reserved



Computer Networking: A Top Down Approach

7th edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks


2.7 socket programming with UDP and TCP

Pure P2P architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

examples:

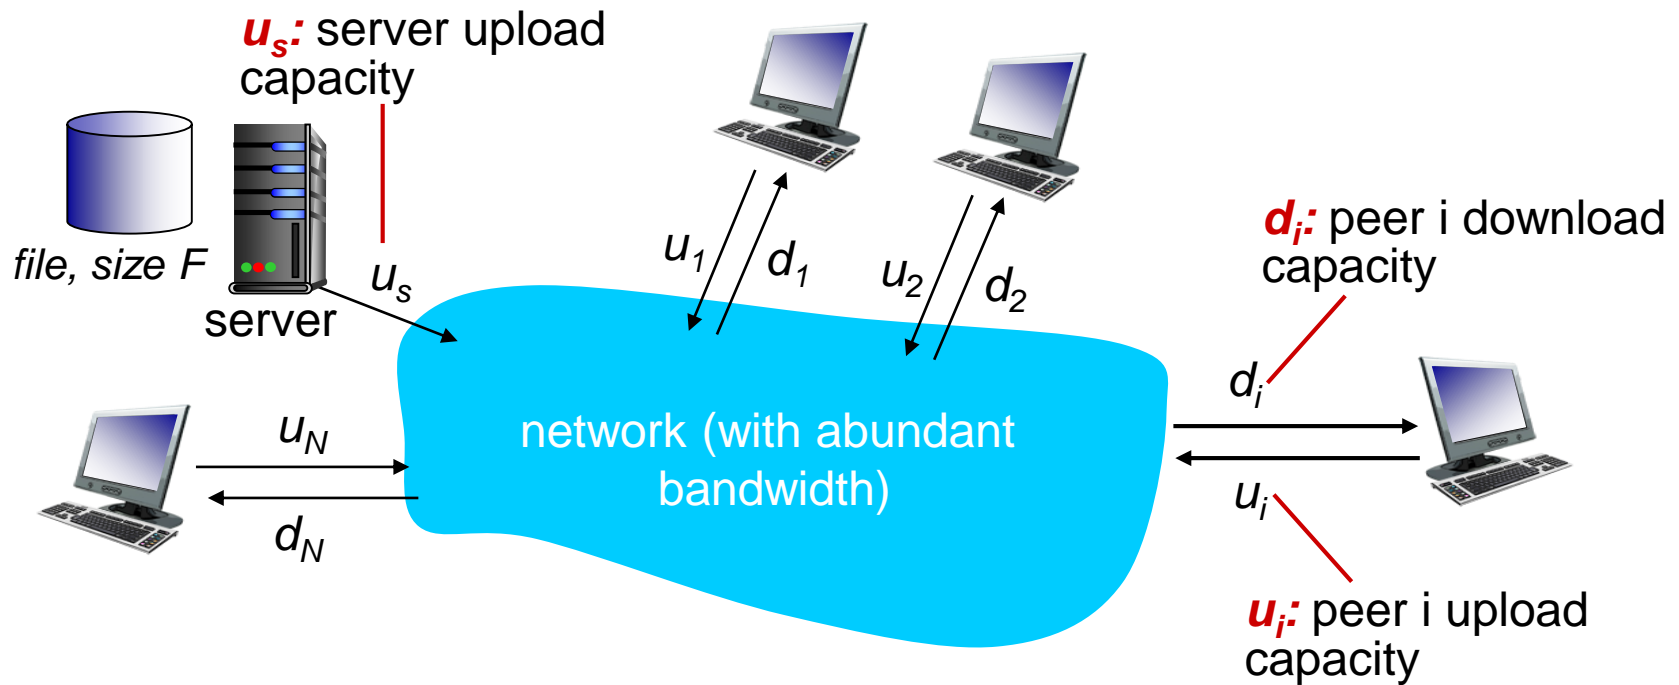
- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)



File distribution: client-server vs P2P

Question: how much time to distribute file (size F) from one server to N peers?

- peer upload/download capacity is limited resource




Side Question

- If a tracker is required, is the tracker legally responsible?
- What if we killed/blocked the tracker?
 - What effect does killing the tracker have?
 - Does it kill the P2P network?
 - Can we somehow still access content?

File distribution time: client-server

- **server transmission:** must sequentially send (upload) N file copies:
 - time to send one copy: F/u_s
 - time to send N copies: NF/u_s
- **client:** each client must download file copy
 - d_{min} = min client download rate
 - min client download time: F/d_{min}




*time to distribute F
to N clients using
client-server approach*

$$D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$$

increases linearly in N

File distribution time: P2P

- **server transmission:** must upload at least one copy
 - time to send one copy: F/u_s
- **client:** each client must download file copy
 - min client download time: F/d_{\min}
- **clients:** as aggregate must download NF bits
 - max upload rate (limiting max download rate) is $u_s + \sum u_i$



*time to distribute F
to N clients using
P2P approach*


$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...

... but so does this, as each peer brings service capacity

Client-server vs. P2P: example

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$




P2P file distribution: BitTorrent

- file divided into 256Kb chunks
- peers in torrent send/receive file chunks

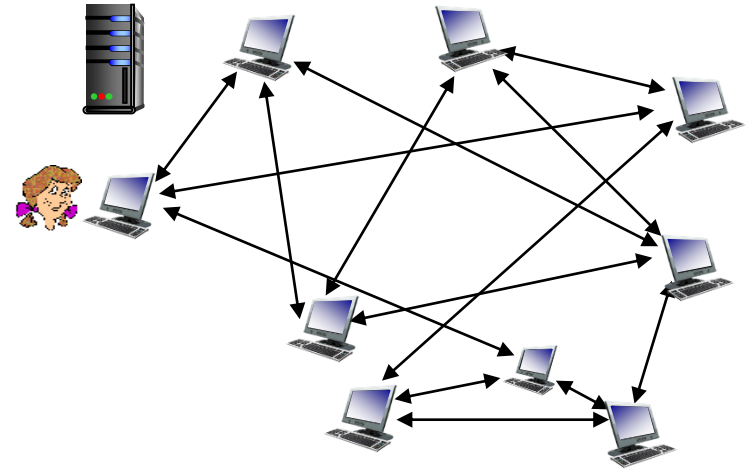
tracker: tracks peers participating in torrent

torrent: group of peers exchanging chunks of a file



P2P file distribution: BitTorrent

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn**: peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



BitTorrent: requesting, sending file chunks

requesting chunks:


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

sending chunks: tit-for-tat

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



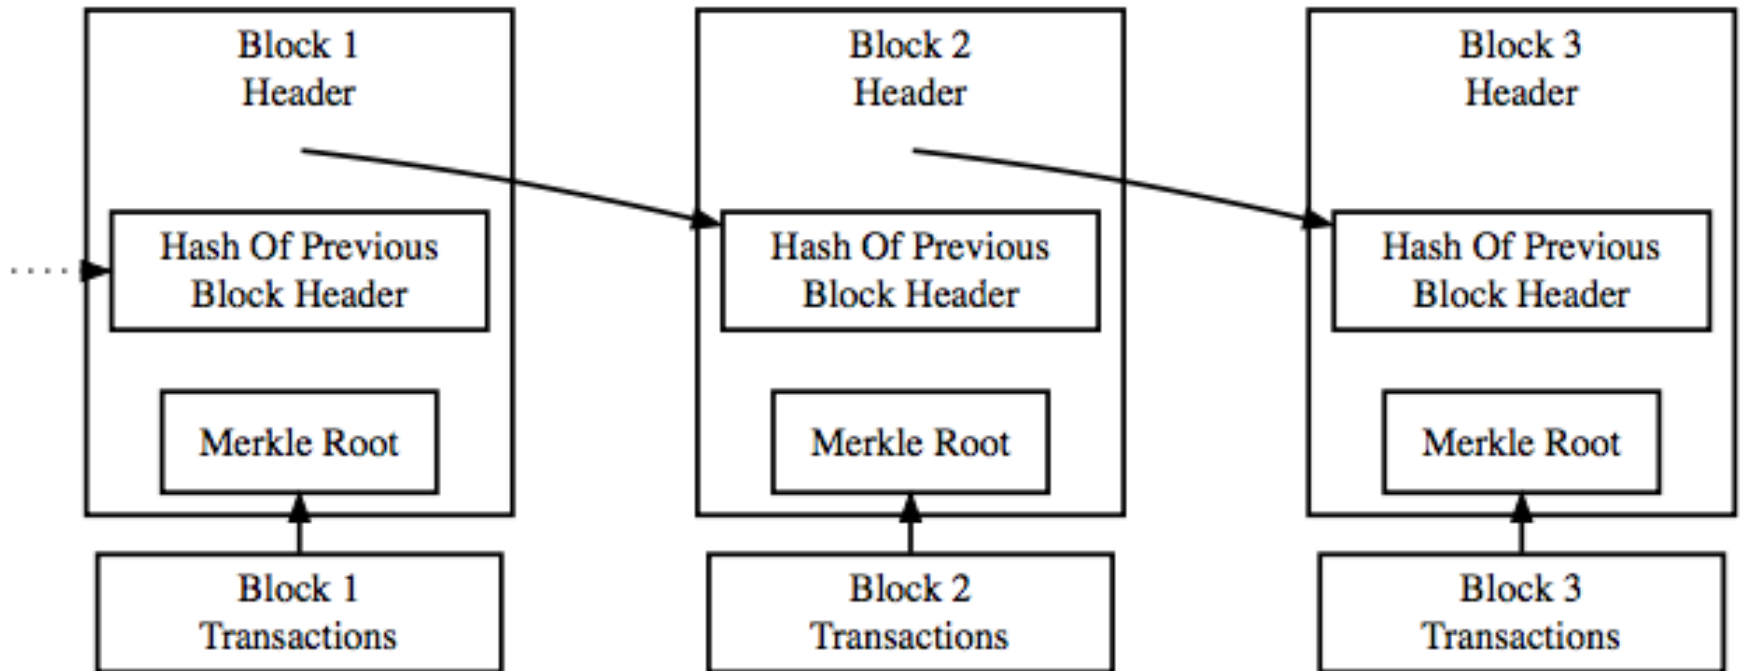
Blockchain

- String of records, or blocks
- Each block contains a hash of the previous block in the chain
- Distributed across many nodes, fairly difficult to modify
 - Nodes are incentivized to reach consensus

Blockchain continued

- Inherently decentralized
 - All nodes have a copy of the blockchain
 - Makes tampering more difficult, must use distributed consensus
- Proof of Work
 - Mining! Finding blocks is hard, verification is easy!
 - Difficulty in mining increases over time

Blockchain illustrated



Simplified Bitcoin Block Chain

Bitcoin

- First widely-adopted cryptocurrency
 - See <https://bitcoin.org/en/developer-guide>
- Built on blockchain transaction ledger
 - Contents within distributed database
 - Bitcoins are awarded for updating/verifying ledger
- Bitcoin are mined
 - Essentially, rewards for perpetuating the network



Kademlia

Network Programming

DHT?

- Why is Kademlia called a Distributed Hash Table?
 - Partition the key space by using node IDs
 - Each ID stores a subset of the table (hence distributed)
 - How do we partition? We use XOR of ID and key as a hash
- So do nodes store <key,value> pairs as a hash table?
 - They can, but don't need to. Often use linked lists.

Understanding Kademlia

- For HW4 we're not making an exact implementation but...
 - [details omitted until HW4 released]
 - Might use SHA-1
 - Might use Python / gRPC
 - Hopefully autograded but scaling will be tricky for the autograder
- Real Kad
 - UDP
 - Viewed as RPCs

Understanding Kademlia

- For HW4 we're not making an exact implementation but...
- Every node generates an ID of length N
 - Fun fact - the behavior for ID collisions is undefined!
- To compute distance between IDs, simply use the bitwise XOR. In Python/C/C++ that's operator^
- Keys are also length N. So we can compute node ID vs key distance the same way.

Node ID Generation

- In real Kad, just pick from 2^{160} randomly
 - If there is a collision, nodes are probably too far to see each other (unless k is large)

OpenSSL SHA-1 Code (C, 1/2)

```
#include <openssl/evp.h>
EVP_MD_CTX *mdctx;
const EVP_MD *md;
char input[] = "12345";
unsigned char id[EVP_MAX_MD_SIZE];
int md_len;

//Create a context
mdctx = EVP_MD_CTX_new();
```

OpenSSL SHA-1 Code (C 2/2)

```
/*Initialize our digest to use the md TYPE object
given by EVP_sha1 */
/*Need to do this every time we want to hash
something */
EVP_DigestInit_ex(mdctx, EVP_sha1(), NULL);
//Add to the hash
EVP_DigestUpdate(mdctx, input, strlen(input));
//Finalize the hash, store in md_value
EVP_DigestFinal_ex(mdctx, id, &md_len);

//Free the context
EVP_MD_CTX_free(mdctx);
```

OpenSSL SHA-1 Code (Python 3)

```
#!/usr/bin/env python3  
import hashlib  
m = hashlib.sha1()  
m.update(b"This is a str")  
m.update(b"ring in two parts")  
m.digest() #Gives the actual hash
```

k and k-buckets

- Routing is split up into k-buckets. There are **N** of them.
- Each bucket $0 \leq i < N$ holds nodes with $2^i \leq \text{distance} < 2^{i+1}$
- If a node is in bucket i , then bit i is the first bit that was different (lowest bit on the left)
- For small i , there are very few possible values, probably an empty bucket.
- For large i , there are many possible values. Don't want to store them all.
- Only store up to k values

k and k-buckets


- How do we pick which ones to kick out?
- Least Recently Used (LRU) list
 - If we follow the [Kad paper](#) the most recently used is the tail and the least recently used is the head
 - Head or tail is fine, as long as it's an LRU list
 - If you receive a message from a node, send it to the "most recent" end of the list
 - We don't want to update LRU when we send to a node, because it might turn out that node is dead. UDP means we won't get a "connection refused"

k and k-buckets

■ Replacement Cache


- We won't do this in HW4 (probably)
- In real Kad, can keep a number of nodes as "backup" - if a spot opens up in k-bucket, use a backup node
- Instead, only kick out if node fails to respond to a PING
- We won't even implement PING, real systems do though it may be very infrequent, e.g. once per 24 hours

Visualizing k-buckets



Source: https://commons.wikimedia.org/wiki/File:Dht_example_SVG.svg

Visualizing k-buckets



Source: <https://pdos.csail.mit.edu/~petar/papers/maymounkov-kademlia-lncs.pdf>

Kademlia Protocol

- Four RPCs

- FIND_NODE(id)

- Returns (IP, UDP Port, Node ID) for k closest nodes to id that are known by the remote node

- FIND_VALUE(key)

- Same thing as FIND_NODE but if the node is storing a value for that key already, then it just returns the value

- PING()

- Check if a node is responding

- STORE(key,value)

- Store the key,value pair. Initiator asks k closest nodes to key.

Doing Lookups

- “Concurrency” parameter α
- Pick up to α nodes from closest non-empty k-bucket (we may only know fewer than α)
- In parallel, make the same RPC (FIND_NODE or FIND_VALUE) to each of those nodes
- RPCs will return nodes, update our buckets and repeat the search with another α nodes (but do not ask nodes we already asked)
- Stop when all k closest nodes have been asked

Kademlia Performance

- Hopefully we can do a search in $O(\log N)$ time, proving this is a little tricky
- $\alpha=1$ behaves similar to Chord for searching performance
- If node IDs are very far away, number of nodes is low, and N and k are not chosen well, it is possible not all nodes can reach each other.